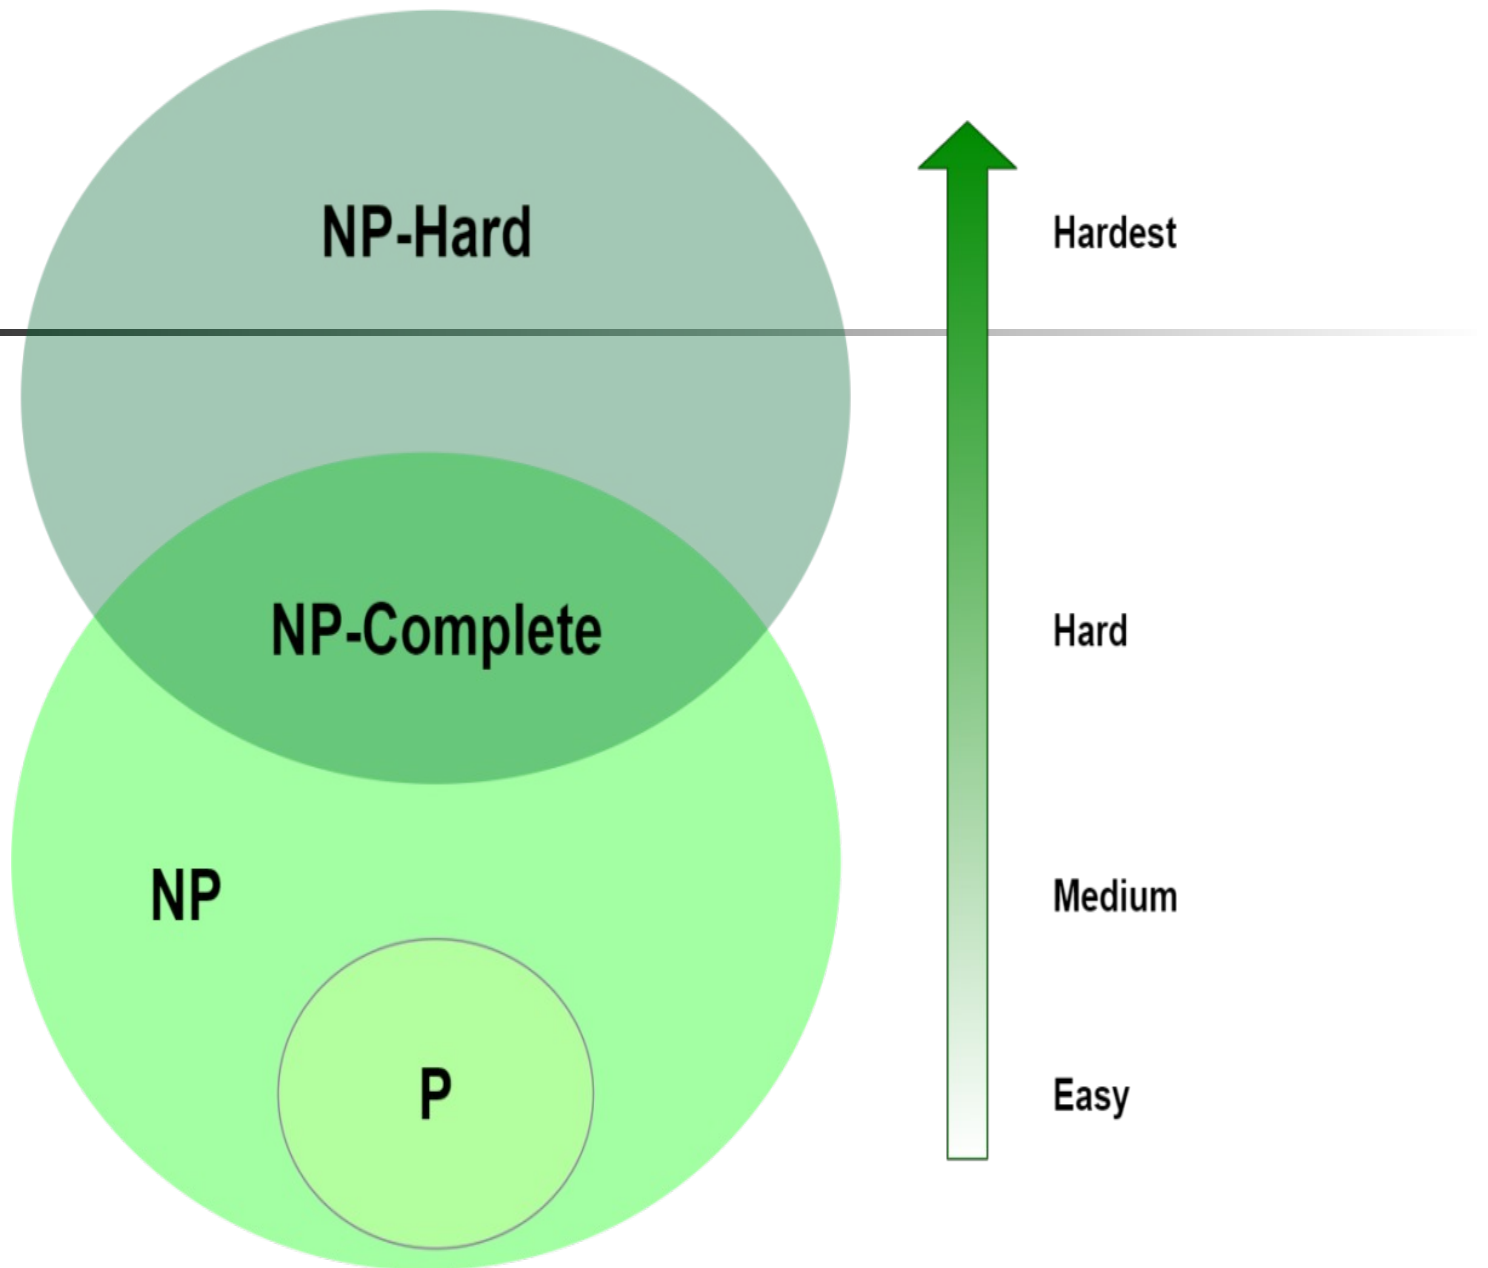





Approximation Algorithms



NP-completeness



“I can’t find an efficient algorithm, but neither can all these famous people.”

Coping With NP-Hardness

Brute-force Algorithms.

- Develop clever enumeration strategies.
- Guaranteed to find optimal solution.
- No guarantees on running time.

Heuristics.

Develop intuitive algorithms.

Guaranteed to run in polynomial time.

No guarantees on quality of solution.

Approximation Algorithms.

- Guaranteed to run in polynomial time.
- Guaranteed to find "high quality" solution, say within 1% of optimum.

Obstacle: need to prove a solution's value is close to optimum, without even knowing what optimum value is!

Coping with NP-completeness

Q. Suppose I need to solve an **NP**-hard optimization problem.
What should I do?

A. Sacrifice one of three desired features.

- i. Runs in polynomial time.
- ii. Solves arbitrary instances of the problem.
- iii. Finds optimal solution to problem.

ρ -approximation algorithm.

- Runs in polynomial time.
- Solves arbitrary instances of the problem
- Finds solution that is within ratio ρ of optimum.

Challenge. Need to prove a solution's value is close to optimum, without even knowing what is optimum value.



Approximation Algorithms

- Up to now, the best algorithm for solving an NP-complete problem requires exponential time in the worst case. It is too time-consuming.
- To reduce the time required for solving a problem, we can relax the problem, and obtain a feasible solution “close” to an optimal solution



Approximation Algorithms

- One compromise is to use **heuristic** solutions.
- The word “heuristic” may be interpreted as “educated guess.”



Approximation Algorithms

An algorithm that returns near-optimal solutions is called an *Approximation Algorithm*.

We need to find an *Approximation Ratio Bound* for an approximation algorithm.



Approximation Ratio Bound

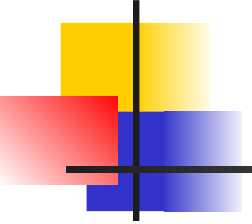
We say an approximation algorithm for the problem has a ratio bound of $\rho(n)$ if for any input size n , the cost C of the solution produced by the approximation algorithm is within a factor of $\rho(n)$ of the C^* of the optimal solution:

$$\max\left\{\frac{C}{C^*}, \frac{C^*}{C}\right\} = \rho(n)$$

This applies for both minimization and maximization problems.

Performance

Guarantees

- 
- An approximation algorithm is bounded by $\rho(n)$ if, for all input of size n , the cost c of the solution obtained by the algorithm is within a factor $\rho(n)$ of the c^* of an optimal solution



ρ -approximation algorithm

- An approximation algorithm with an approximation ratio bound of ρ is called a ρ -approximation algorithm or a $(1+\varepsilon)$ -approximation algorithm.
- Note that ρ is always larger than 1 and $\varepsilon = \rho - 1$.







Vertex Cover Problem

- Let $G=(V, E)$. The subset S of V that meets every edge of E is referred to as the **Vertex Cover**.
- The Vertex Cover Problem is solved for finding a vertex cover of the **Minimum** size. It is NP-hard Computational Problem or the Optimization Version of an NP-Complete Decision Problem.



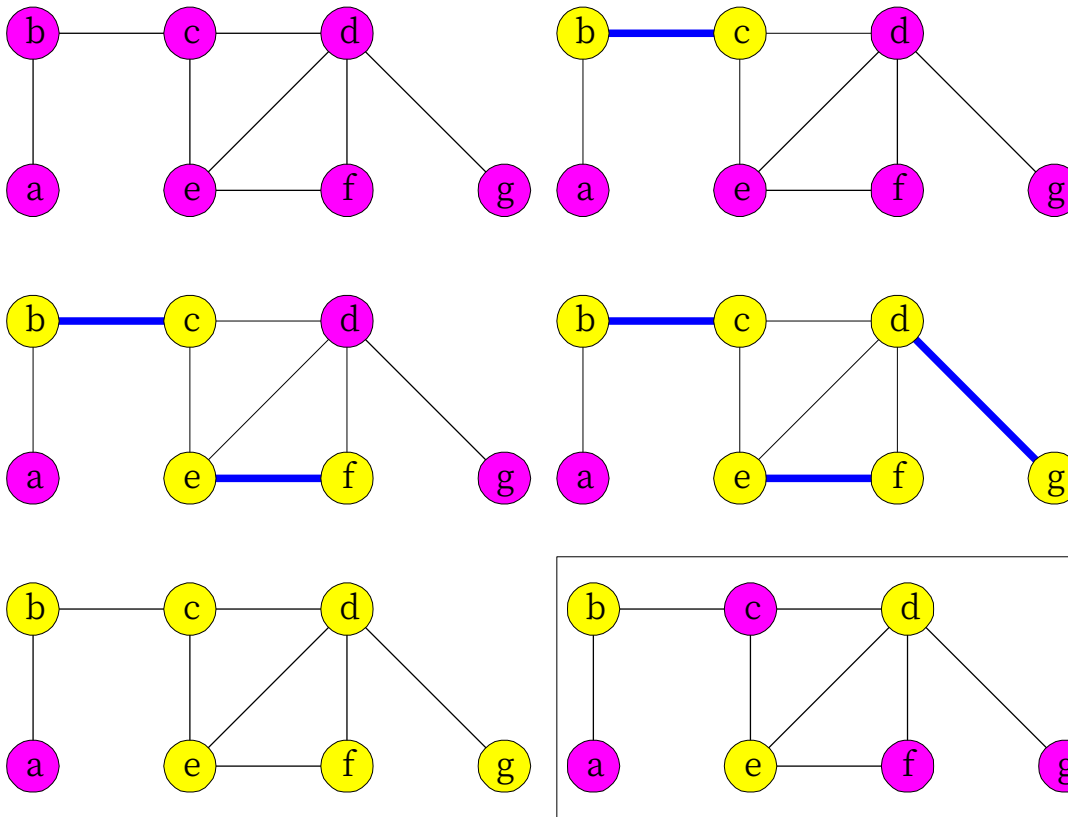
Examples of Vertex Cover



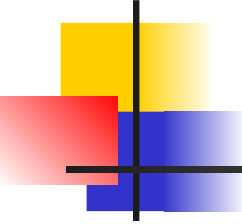


APPROX_VERTEX_COVER(G)

```
1   $C \leftarrow \phi$ 
2   $E' \leftarrow E(G)$ 
3  while  $E' \neq \phi$ 
4      do let  $(u, v)$  be an arbitrary edge of  $E'$ 
5           $C \leftarrow C \cup \{u, v\}$ 
6          remove from  $E'$  every edge incident on either  $u$  or  $v$ 
7  return  $C$ 
```



Complexity: $O(E)$



Theorem: APPROX_VERTEX_COVER has ratio bound of 2.

Proof.

C^* : optimal solution

C : approximate solution

A : the set of edges selected in

step 4

Let A be the set of selected edges.

$$|C| = 2|A|$$

$$|A| \leq |C^*|$$

When one edge is selected, 2 vertices are added into C .

No two edges in A share a common endpoint.

$$\Rightarrow |C| \leq 2|C^*|$$




Bin Packing Problem

- Given n items of sizes a_1, a_2, \dots, a_n , $0 < a_i \leq 1$ for $1 \leq i \leq n$, which have to be placed in bins of unit capability, the bin packing problem is solved for determining the minimum number of bins to accommodate all items.
- If we consider the items of different sizes to be the lengths of time of executing different jobs on a standard processor, then the problem becomes to use minimum number of processors which can finish all of the jobs within a fixed time. // We can assume the longest job takes one unit time, which equals

Example of Bin Packing Problem

- Ex. Given $n = 5$ items with sizes 0.3, 0.5, 0.8, 0.2, 0.4, the optimal solution is 3 bins.



The bin packing problem is NP-hard optimization problem.

An Approximation Algorithm for the Bin Packing Problem



- An Approximation Algorithm: (First-Fit (FF)) place the item i into the lowest-indexed bin which can accommodate the item i .
- OPT: The number of bins of the Optimal Solution
- FF: The number of bins in the First-Fit Algorithm
- $C(B_i)$: The sum of the sizes of items packed in bin B_i in the First-Fit Algorithm
- Let $FF=m$.

An Approximation Algorithm for the Bin Packing Problem

- $OPT \geq \left\lceil \sum_{i=1}^n a_i \right\rceil$, ceiling of sum of sizes of all items
- $C(B_i) + C(B_{i+1}) > 1$ (b) (Otherwise, the items in B_{i+1} will be put in B_i).
- $C(B_1) + C(B_m) > 1$ (b) (Otherwise, the items in B_m will be put in B_1 .)
- For m nonempty bins, $\sum_{i=1}^m C(B_i) > \sum_{i=1}^n a_i$, (a)+(b) for $i=1, \dots, m$
 $\Rightarrow FF = m < 2 \sum_{i=1}^n a_i = 2 \leq 2 OPT$

Load balancing

Input. m identical machines; $n \geq m$ jobs, job j has processing time t_j .

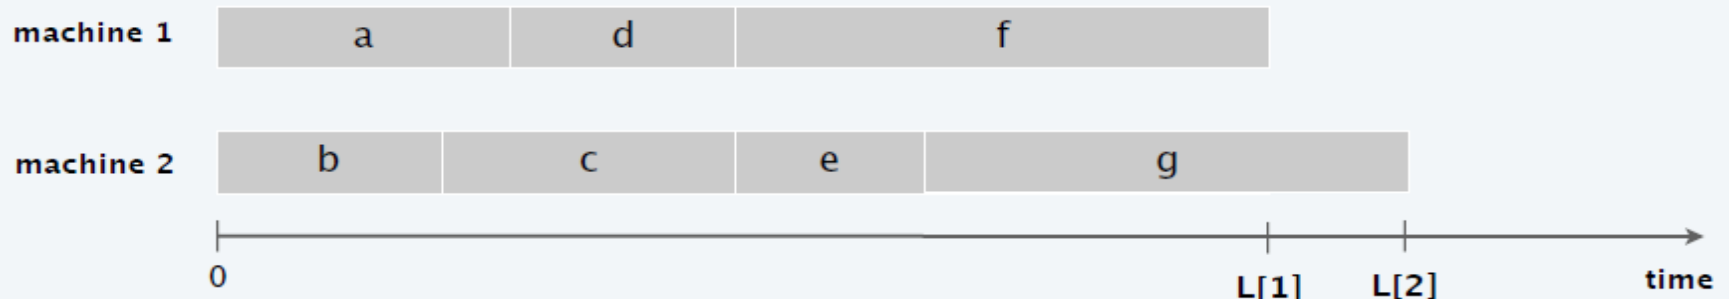
- Job j must run contiguously on one machine.
- A machine can process at most one job at a time.

Def. Let $S[i]$ be the subset of jobs assigned to machine i .

The **load** of machine i is $L[i] = \sum_{j \in S[i]} t_j$.

Def. The **makespan** is the maximum load on any machine $L = \max_i L[i]$.

Load balancing. Assign each job to a machine to minimize makespan.

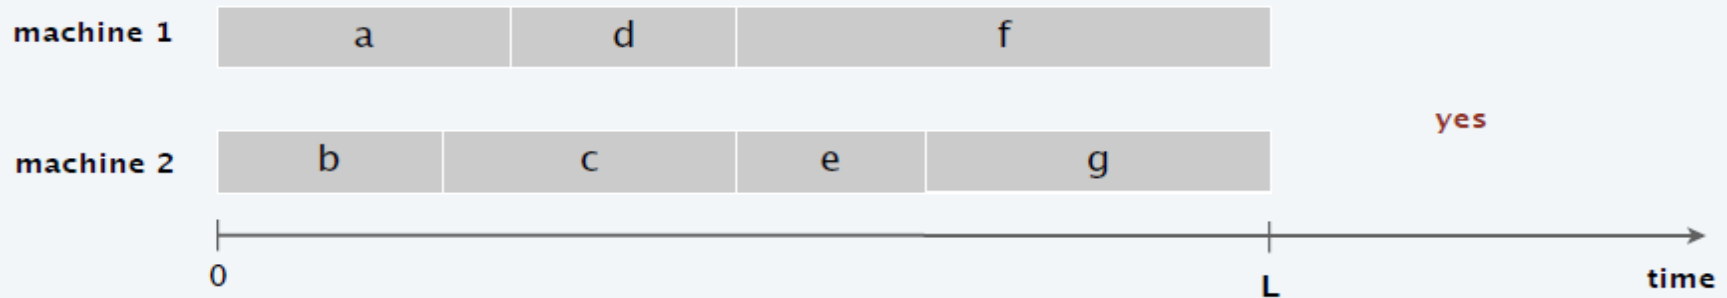
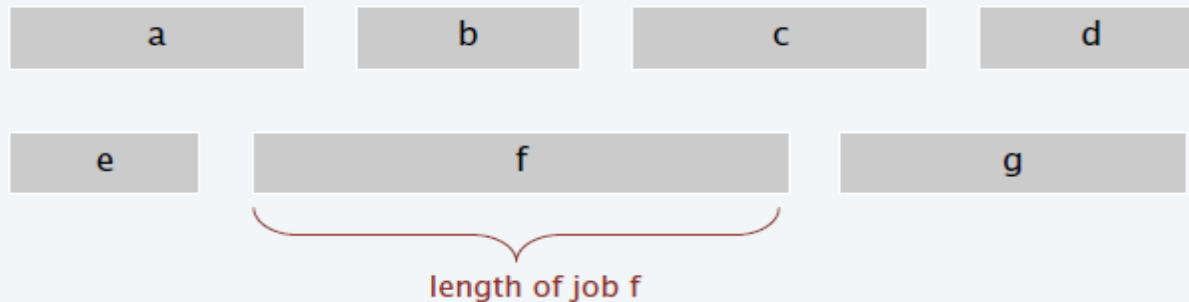


Load balancing on 2 machines is NP-hard

Claim. Load balancing is hard even if $m = 2$ machines.

Pf. $\text{PARTITION} \leq_p \text{LOAD-BALANCE}$.

NP-complete by Exercise 8.26



Load balancing: list scheduling

List-scheduling algorithm.

- Consider n jobs in some fixed order.
- Assign job j to machine i whose load is smallest so far.

LIST-SCHEDULING ($m, n, t_1, t_2, \dots, t_n$)

FOR $i = 1$ TO m

$L[i] \leftarrow 0.$ \leftarrow load on machine i

$S[i] \leftarrow \emptyset.$ \leftarrow jobs assigned to machine i

FOR $j = 1$ TO n

$i \leftarrow \operatorname{argmin}_k L[k].$ \leftarrow machine i has smallest load

$S[i] \leftarrow S[i] \cup \{j\}.$ \leftarrow assign job j to machine i

$L[i] \leftarrow L[i] + t_j.$ \leftarrow update load of machine i

RETURN $S[1], S[2], \dots, S[m].$

Implementation. $O(n \log m)$ using a priority queue for loads $L[k]$.

Load balancing: list scheduling analysis

Theorem. [Graham 1966] Greedy algorithm is a 2-approximation.

- First worst-case analysis of an approximation algorithm.
- Need to compare resulting solution with optimal makespan L^* .

Lemma 1. For all k : the optimal makespan $L^* \geq t_k$.

Pf. Some machine must process the most time-consuming job. ■

Lemma 2. The optimal makespan $L^* \geq \frac{1}{m} \sum_k t_k$.

Pf.

- The total processing time is $\sum_k t_k$.
- One of m machines must do at least a $1 / m$ fraction of total work. ■


Load balancing: list scheduling analysis

Theorem. Greedy algorithm is a 2-approximation.

Pf. Consider load $L[i]$ of bottleneck machine i . ← machine that ends up with highest load


- Let j be last job scheduled on machine i .
- When job j assigned to machine i , i had smallest load.

Its load before assignment is $L[i] - t_j$; hence $L[i] - t_j \leq L[k]$ for all $1 \leq k \leq m$.



Load balancing: list scheduling analysis

Theorem. Greedy algorithm is a 2-approximation.

Pf. Consider load $L[i]$ of bottleneck machine i .  machine that ends up with highest load

- Let j be last job scheduled on machine i .
- When job j assigned to machine i , i had smallest load.

Its load before assignment is $L[i] - t_j$; hence $L[i] - t_j \leq L[k]$ for all $1 \leq k \leq m$.

- Sum inequalities over all k and divide by m :

$$\begin{aligned} L[i] - t_j &\leq \frac{1}{m} \sum_k L[k] \\ &= \frac{1}{m} \sum_k t_k \end{aligned}$$

Lemma 2  $\leq L^*$.

- Now, $L = L[i] = (L[i] - t_j) + t_j \leq 2L^*$.

$$\begin{array}{ccc} \underbrace{}_{\leq L^*} & + & \underbrace{}_{\leq L^*} \\ \uparrow & & \uparrow \\ \text{above inequality} & & \text{Lemma 1} \end{array}$$

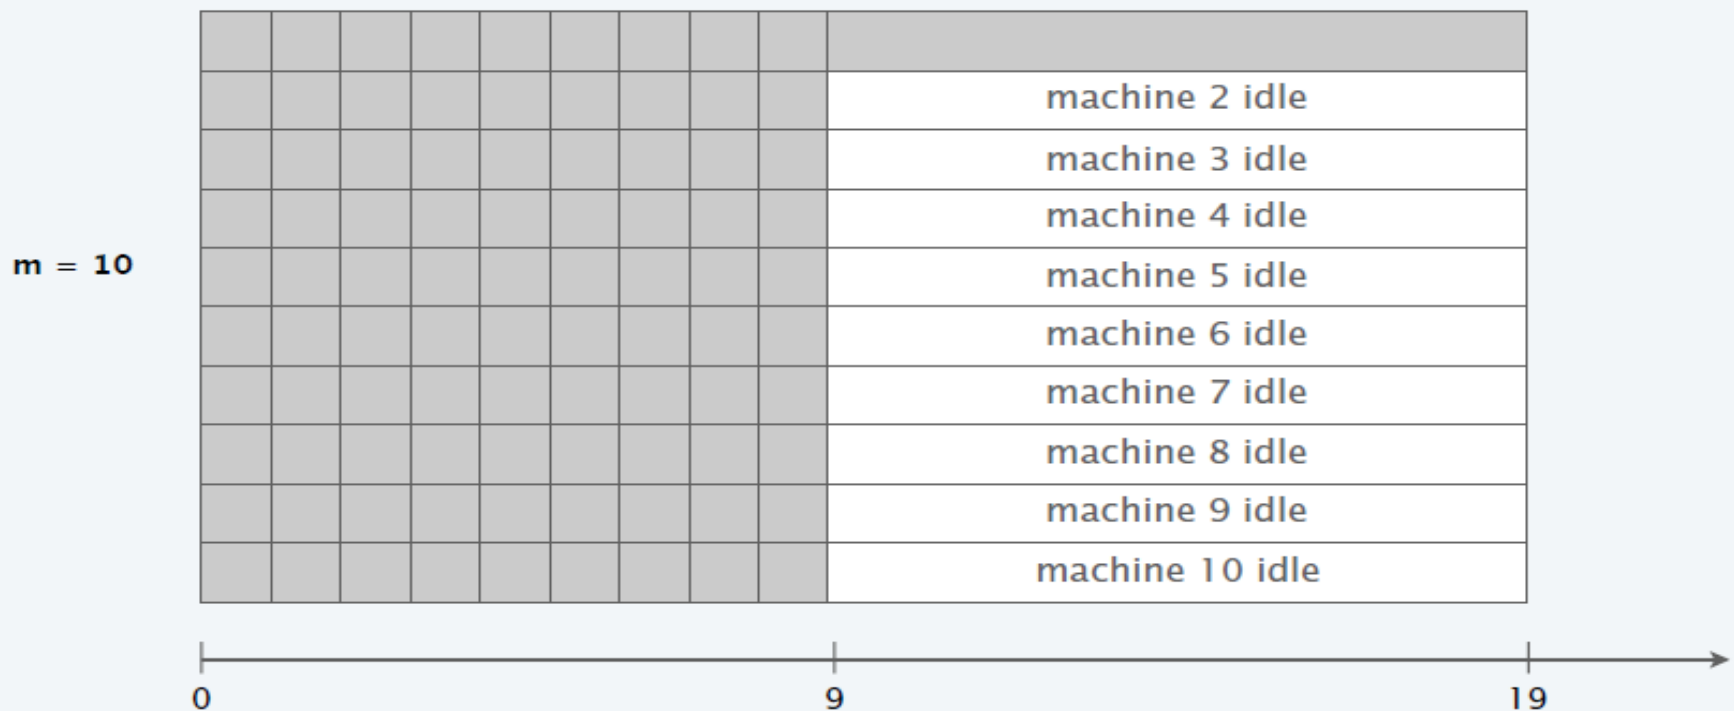
Load balancing: list scheduling analysis

Q. Is our analysis tight?

A. Essentially yes.

Ex: m machines, first $m(m-1)$ jobs have length 1, last job has length m .

list scheduling makespan = $19 = 2m - 1$



Load balancing: list scheduling analysis


Q. Is our analysis tight?

A. Essentially yes.

Ex: m machines, first $m(m-1)$ jobs have length 1, last job has length m .

optimal makespan = $10 = m$

$m = 10$



Load balancing: LPT rule

Longest processing time (LPT). Sort n jobs in decreasing order of processing times; then run list scheduling algorithm.

LPT-LIST-SCHEDULING ($m, n, t_1, t_2, \dots, t_n$)

SORT jobs and renumber so that $t_1 \geq t_2 \geq \dots \geq t_n$.

FOR $i = 1$ TO m

$L[i] \leftarrow 0.$ \leftarrow load on machine i

$S[i] \leftarrow \emptyset.$ \leftarrow jobs assigned to machine i

FOR $j = 1$ TO n

$i \leftarrow \operatorname{argmin}_k L[k].$ \leftarrow machine i has smallest load

$S[i] \leftarrow S[i] \cup \{j\}.$ \leftarrow assign job j to machine i

$L[i] \leftarrow L[i] + t_j.$ \leftarrow update load of machine i

RETURN $S[1], S[2], \dots, S[m].$

Load balancing: LPT rule

Observation. If bottleneck machine i has only 1 job, then optimal.

Pf. Any solution must schedule that job. ■

Lemma 3. If there are more than m jobs, $L^* \geq 2t_{m+1}$.

Pf.

- Consider processing times of first $m+1$ jobs $t_1 \geq t_2 \geq \dots \geq t_{m+1}$.
- Each takes at least t_{m+1} time.
- There are $m+1$ jobs and m machines, so by pigeonhole principle, at least one machine gets two jobs. ■

Theorem. LPT rule is a $3/2$ -approximation algorithm.

Pf. [similar to proof for list scheduling]

- Consider load $L[i]$ of bottleneck machine i .
- Let j be last job scheduled on machine i . ← assuming machine i has at least 2 jobs, we have $j \geq m+1$

$$L = L[i] = \underbrace{(L[i] - t_j)}_{\text{as before}} + \underbrace{t_j}_{\leq \frac{1}{2} L^*} \leq \frac{3}{2} L^* \quad \blacksquare$$

as before $\rightarrow \leq L^* \quad \leq \frac{1}{2} L^* \leftarrow \text{Lemma 3 (since } t_{m+1} \geq t_j)$

Load balancing: LPT rule

Q. Is our $3/2$ analysis tight?

A. No.

Theorem. [Graham 1969] LPT rule is a $4/3$ -approximation.

Pf. More sophisticated analysis of same algorithm.

Q. Is Graham's $4/3$ analysis tight?

A. Essentially yes.

Ex.

- m machines
- $n = 2m + 1$ jobs
- 2 jobs of length $m, m + 1, \dots, 2m - 1$ and one more job of length m .
- Then, $L / L^* = (4m - 1) / (3m)$

Generalized load balancing

Input. Set of m machines M ; set of n jobs J .

- Job $j \in J$ must run contiguously on an **authorized machine** in $M_j \subseteq M$.
- Job $j \in J$ has processing time t_j .
- Each machine can process at most one job at a time.

Def. Let J_i be the subset of jobs assigned to machine i .

The load of machine i is $L_i = \sum_{j \in J_i} t_j$.

Def. The makespan is the maximum load on any machine $= \max_i L_i$.

Generalized load balancing. Assign each job to an authorized machine to minimize makespan.

Generalized load balancing: integer linear program and relaxation

ILP formulation. x_{ij} = time machine i spends processing job j .

$$\begin{aligned} (IP) \quad & \min \quad L \\ & \text{s. t.} \quad \sum_i x_{ij} = t_j \quad \text{for all } j \in J \\ & \quad \sum_j x_{ij} \leq L \quad \text{for all } i \in M \\ & \quad x_{ij} \in \{0, t_j\} \quad \text{for all } j \in J \text{ and } i \in M_j \\ & \quad x_{ij} = 0 \quad \text{for all } j \in J \text{ and } i \notin M_j \end{aligned}$$

LP relaxation.

$$\begin{aligned} (LP) \quad & \min \quad L \\ & \text{s. t.} \quad \sum_i x_{ij} = t_j \quad \text{for all } j \in J \\ & \quad \sum_j x_{ij} \leq L \quad \text{for all } i \in M \\ & \quad x_{ij} \geq 0 \quad \text{for all } j \in J \text{ and } i \in M_j \\ & \quad x_{ij} = 0 \quad \text{for all } j \in J \text{ and } i \notin M_j \end{aligned}$$

Generalized load balancing: lower bounds

Lemma 1. The optimal makespan $L^* \geq \max_j t_j$.

Pf. Some machine must process the most time-consuming job. ■

Lemma 2. Let L be optimal value to the LP . Then, optimal makespan $L^* \geq L$.


Pf. LP has fewer constraints than ILP formulation. ■

Generalized load balancing: structure of LP solution


Lemma 3. Let x be solution to LP . Let $G(x)$ be the graph with an edge between machine i and job j if $x_{ij} > 0$. Then $G(x)$ is **acyclic**.

Pf. (deferred)

can transform x into another LP solution where $G(x)$ is acyclic if LP solver doesn't return such an x



$G(x)$ acyclic



$x_{ij} > 0$

$G(x)$ cyclic

○ job

□ machine

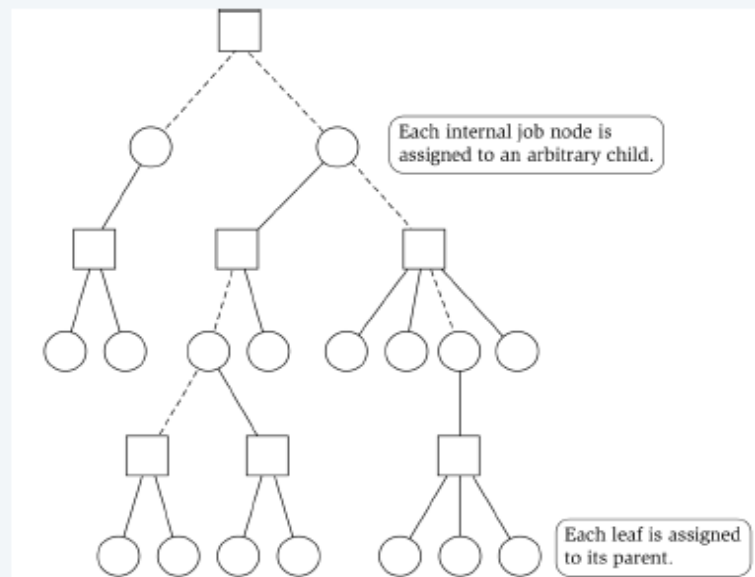

Generalized load balancing: rounding

Rounded solution. Find LP solution x where $G(x)$ is a forest. Root forest $G(x)$ at some arbitrary machine node r .

- If job j is a leaf node, assign j to its parent machine i .
- If job j is not a leaf node, assign j to any one of its children.

Lemma 4. Rounded solution only assigns jobs to authorized machines.

Pf. If job j is assigned to machine i , then $x_{ij} > 0$. LP solution can only assign positive value to authorized machines. ■



Generalized load balancing: analysis

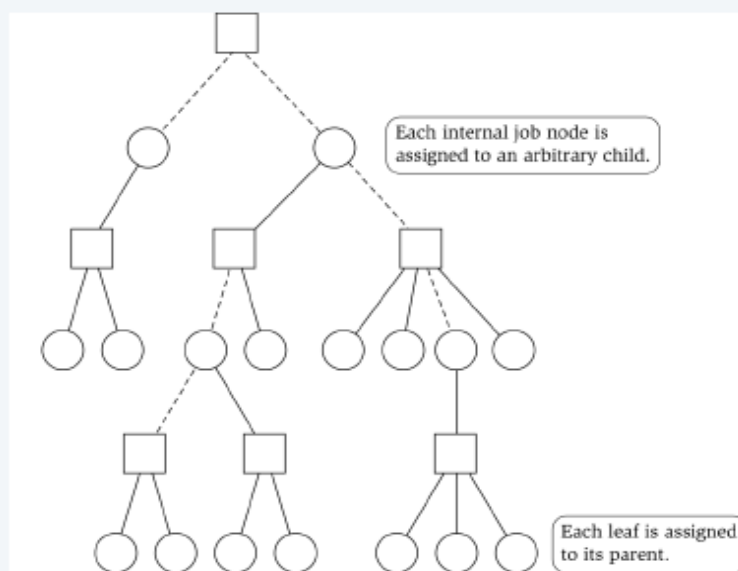

Lemma 5. If job j is a leaf node and machine $i = \text{parent}(j)$, then $x_{ij} = t_j$.

Pf.

- Since i is a leaf, $x_{ij} = 0$ for all $j \neq \text{parent}(i)$.
- LP constraint guarantees $\sum_i x_{ij} = t_j$. ■

Lemma 6. At most one non-leaf job is assigned to a machine.

Pf. The only possible non-leaf job assigned to machine i is $\text{parent}(i)$. ■



Generalized load balancing: analysis

Theorem. Rounded solution is a 2-approximation.

Pf.

- Let $J(i)$ be the jobs assigned to machine i .
- By LEMMA 6, the load L_i on machine i has two components:

- leaf nodes:

$$\sum_{\substack{j \in J(i) \\ j \text{ is a leaf}}} t_j \stackrel{\text{Lemma 5}}{=} \sum_{\substack{j \in J(i) \\ j \text{ is a leaf}}} x_{ij} \leq \sum_{j \in J} x_{ij} \leq L \stackrel{\text{LP}}{\leq} L^* \stackrel{\text{Lemma 2 (LP is a relaxation)}}{\leq} L^*$$

optimal value of LP

- parent:

$$t_{\text{parent}(i)} \stackrel{\text{Lemma 1}}{\leq} L^*$$

- Thus, the overall load $L_i \leq 2L^*$. ■

Generalized load balancing: flow formulation


Flow formulation of LP .

$$\sum_i x_{ij} = t_j \quad \text{for all } j \in J$$

$$\sum_j x_{ij} \leq L \quad \text{for all } i \in M$$

$$x_{ij} \geq 0 \quad \text{for all } j \in J \text{ and } i \in M_j$$

$$x_{ij} = 0 \quad \text{for all } j \in J \text{ and } i \notin M_j$$





Observation. Solution to feasible flow problem with value L are in 1-to-1 correspondence with LP solutions of value L .

Generalized load balancing: structure of solution

Lemma 3. Let (x, L) be solution to LP . Let $G(x)$ be the graph with an edge from machine i to job j if $x_{ij} > 0$. We can find another solution (x', L) such that $G(x')$ is acyclic.

Pf. Let C be a cycle in $G(x)$.

- Augment flow along the cycle C . \leftarrow flow conservation maintained
- At least one edge from C is removed (and none are added).
- Repeat until $G(x')$ is acyclic. ■



Conclusions

Running time. The bottleneck operation in our 2-approximation is solving one LP with $mn + 1$ variables.

Remark. Can solve LP using flow techniques on a graph with $m + n + 1$ nodes: given L , find feasible flow if it exists. Binary search to find L^* .

Extensions: unrelated parallel machines. [Lenstra–Shmoys–Tardos 1990]

- Job j takes t_{ij} time if processed on machine i .
- 2-approximation algorithm via LP rounding.
- If $P \neq NP$, then no ρ -approximation exists for any $\rho < 3/2$.