



Department of Information Technology
National Institute of Technology Karnataka, Surathkal

Distributed Memory Parallelism with MPI


**By,
Geetha V
Dept of IT,
NITK Surathkal**

Outline

- **Distributed Memory Architecture**
- **Introduction to MPI**
- **Structure of MPI program**
- **Types of Message Passing**
- **Basic Routines in Point to Point Communication**
- **Example programs on Point to Point Communication**
- **Basic Routines in Collective Communication**
- **Sample Programs on Collective Communication**

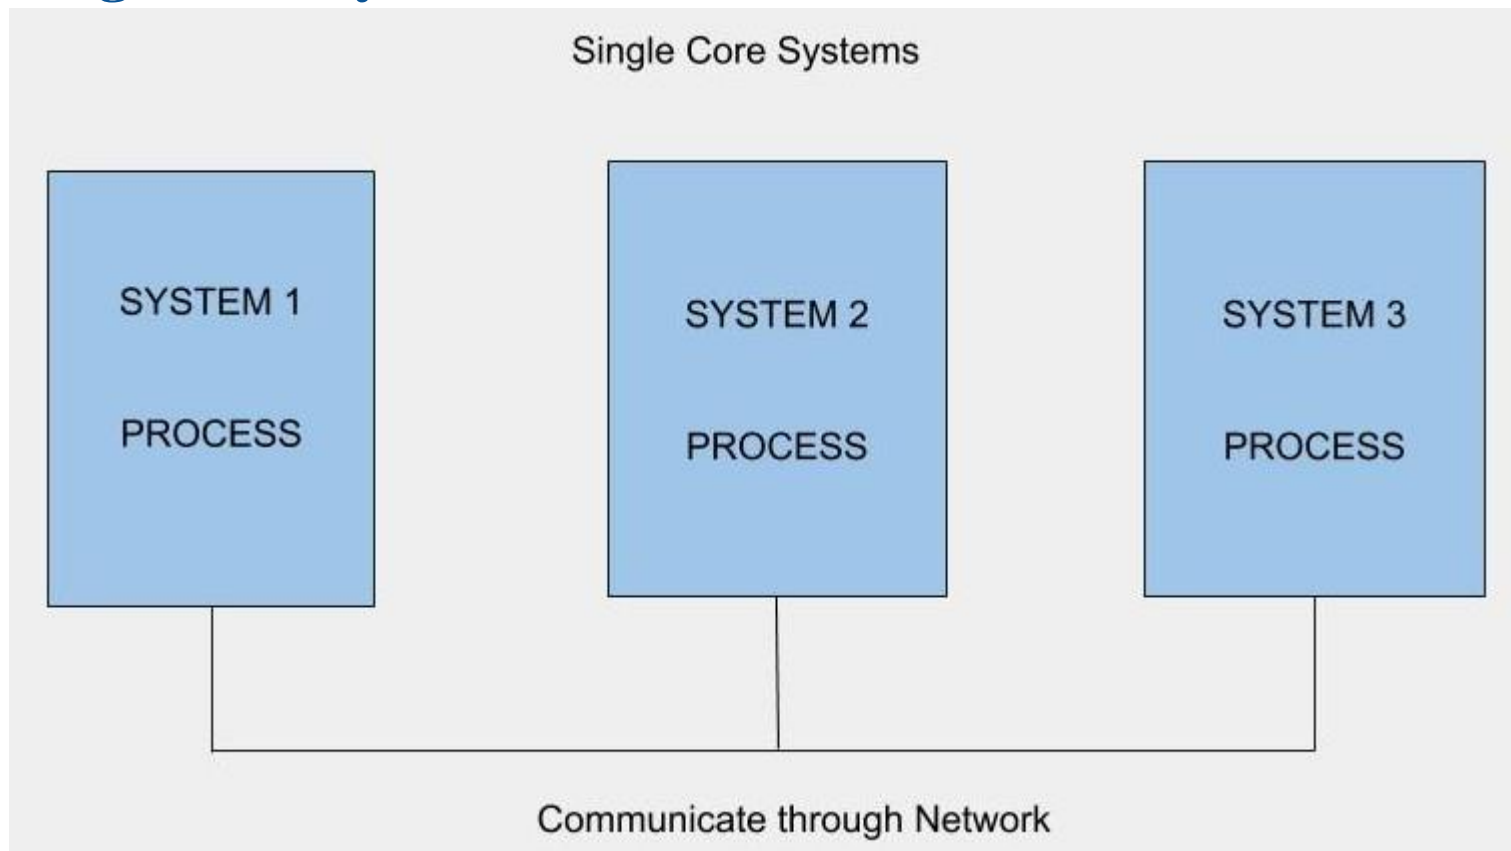
Distributed Memory Architecture

- Each processor has its own memory
- They cannot access the memory of other processors.
- Any data that needs to be shared must be explicitly transmitted from one processor to another using Message Passing.



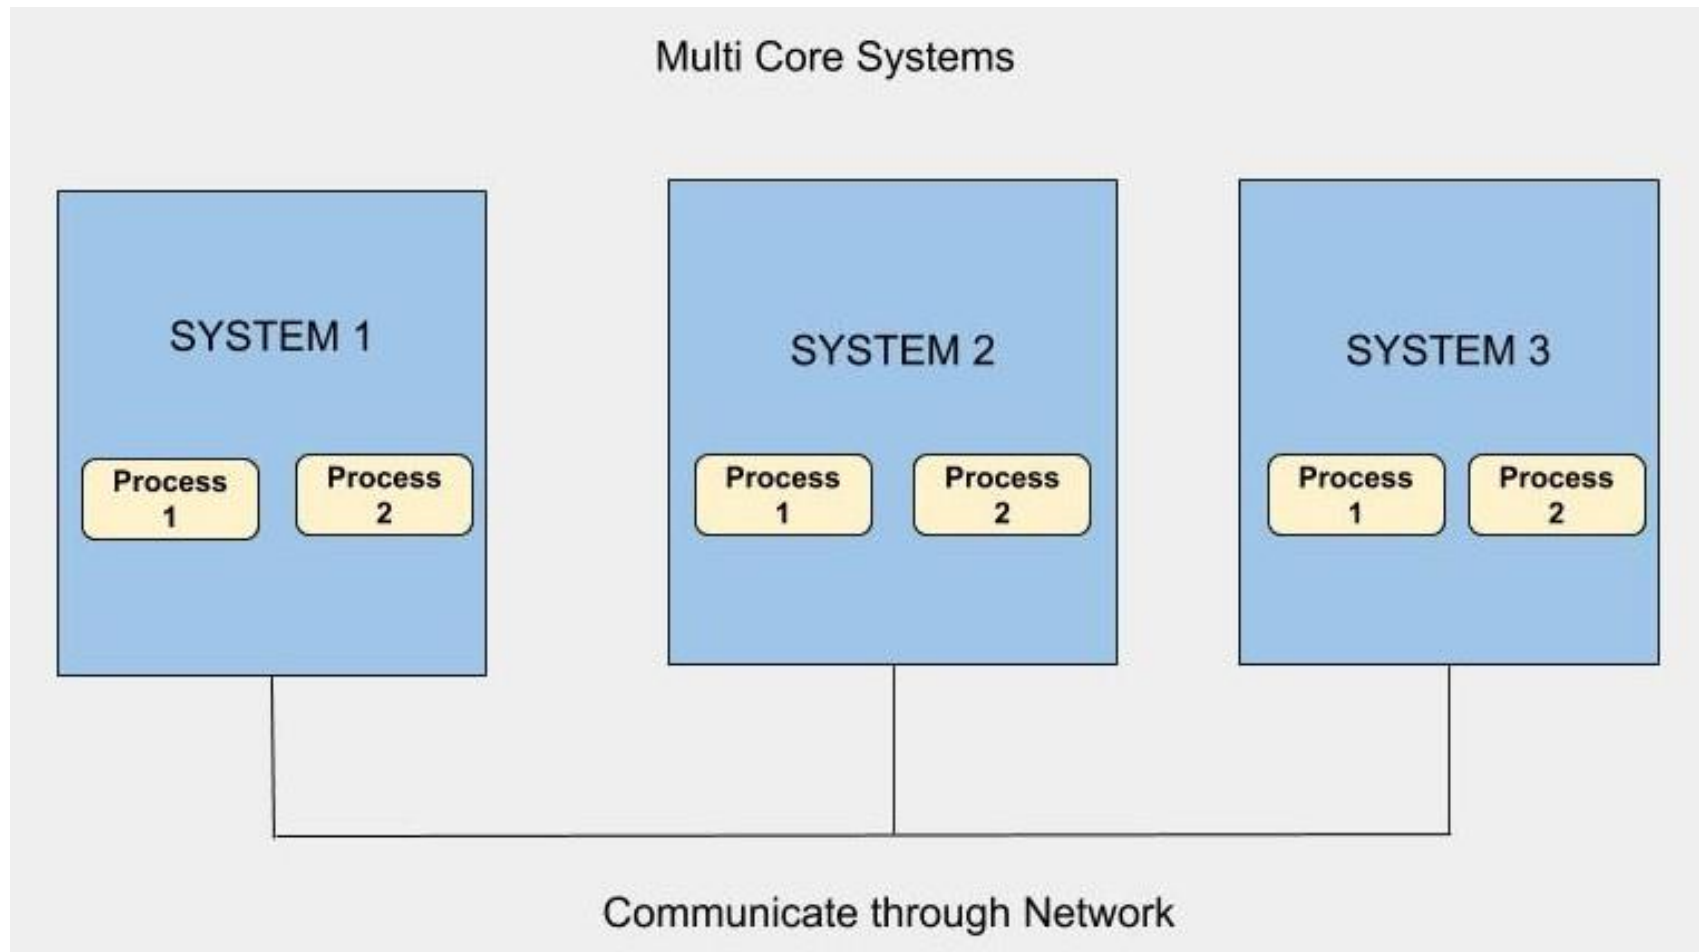
DISTRIBUTED MEMORY ARCHITECTURE

- Systems with single core communicating through distributed memory.
- Heterogeneous systems




DISTRIBUTED MEMORY ARCHITECTURE

- Systems with multiple core communicating through shared and distributed memory




Hybrid Model



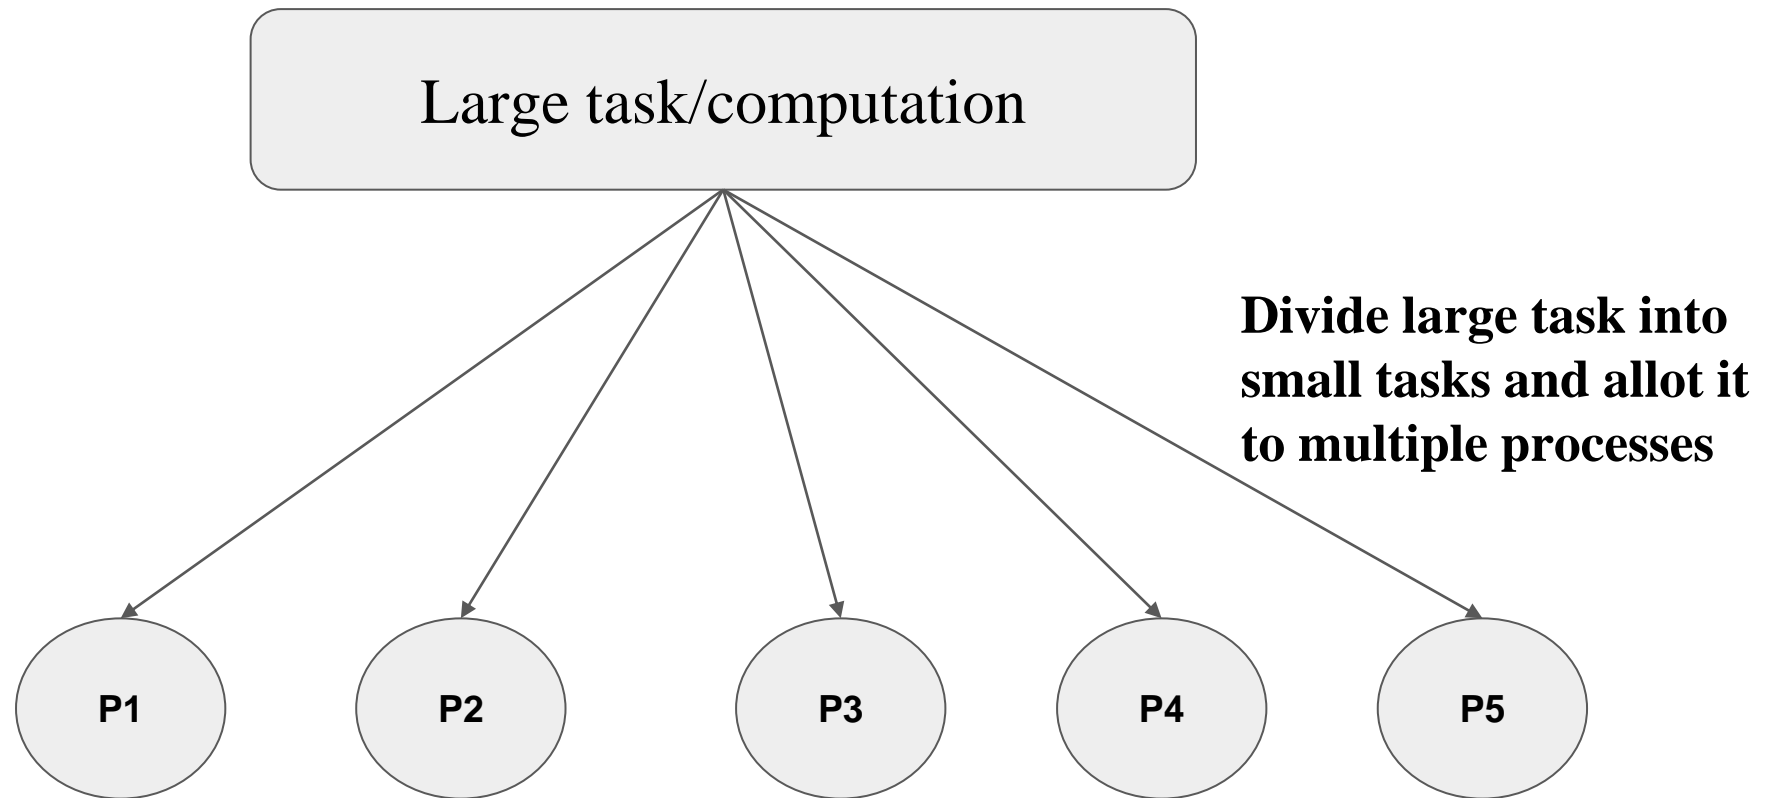
Reference: https://computing.llnl.gov/tutorials/parallel_comp/#ModelsMessage

Hybrid Model



Reference: https://computing.llnl.gov/tutorials/parallel_comp/#ModelsMessage

Parallel Computation:



For Example: $N = 1,00,000$ divided into $P1=20,000$, $P2=20,000$

Computation is same. Data is different. Single Program Multiple Data

INTRODUCTION TO MPI

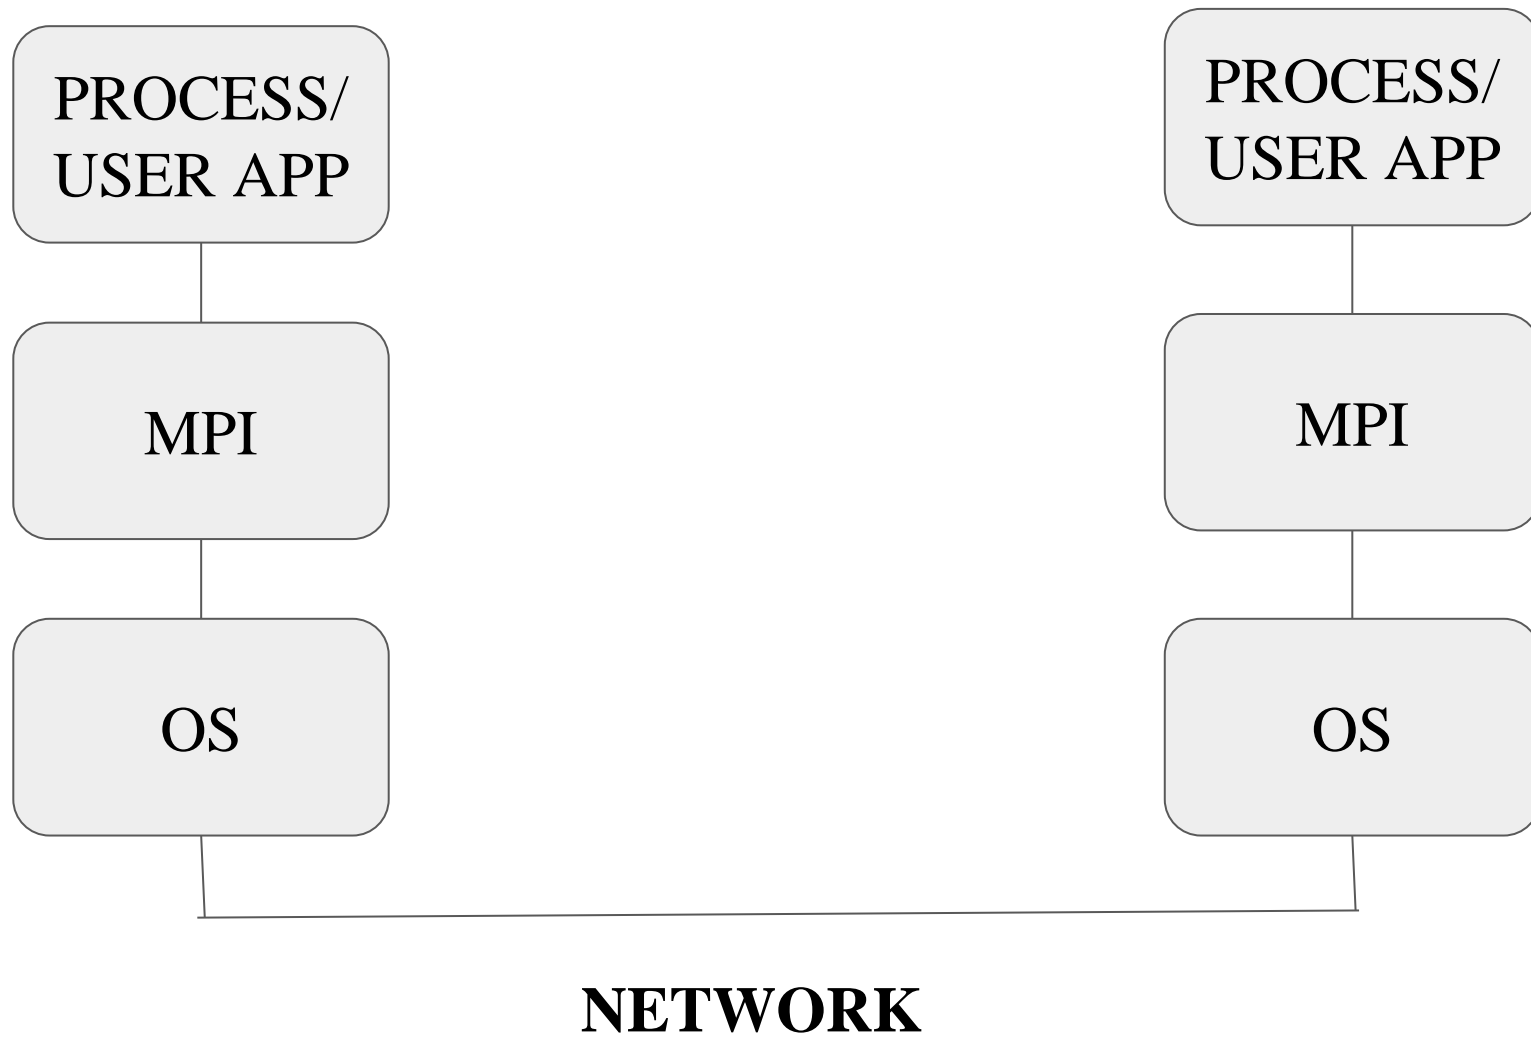
- **What is MPI?**

- Message Passing Interface is a specification.
 - A standard for vendors to implement.
- It is a library, i.e. a set of subroutines, functions and constants
- Allows Message Passing between processes.
- It is based on Single Program, Multiple Data (SPMD)
 - Every process executes the same program
 - Each process performs computations on its local variables, then communicates with other processes, in order to get the final result.

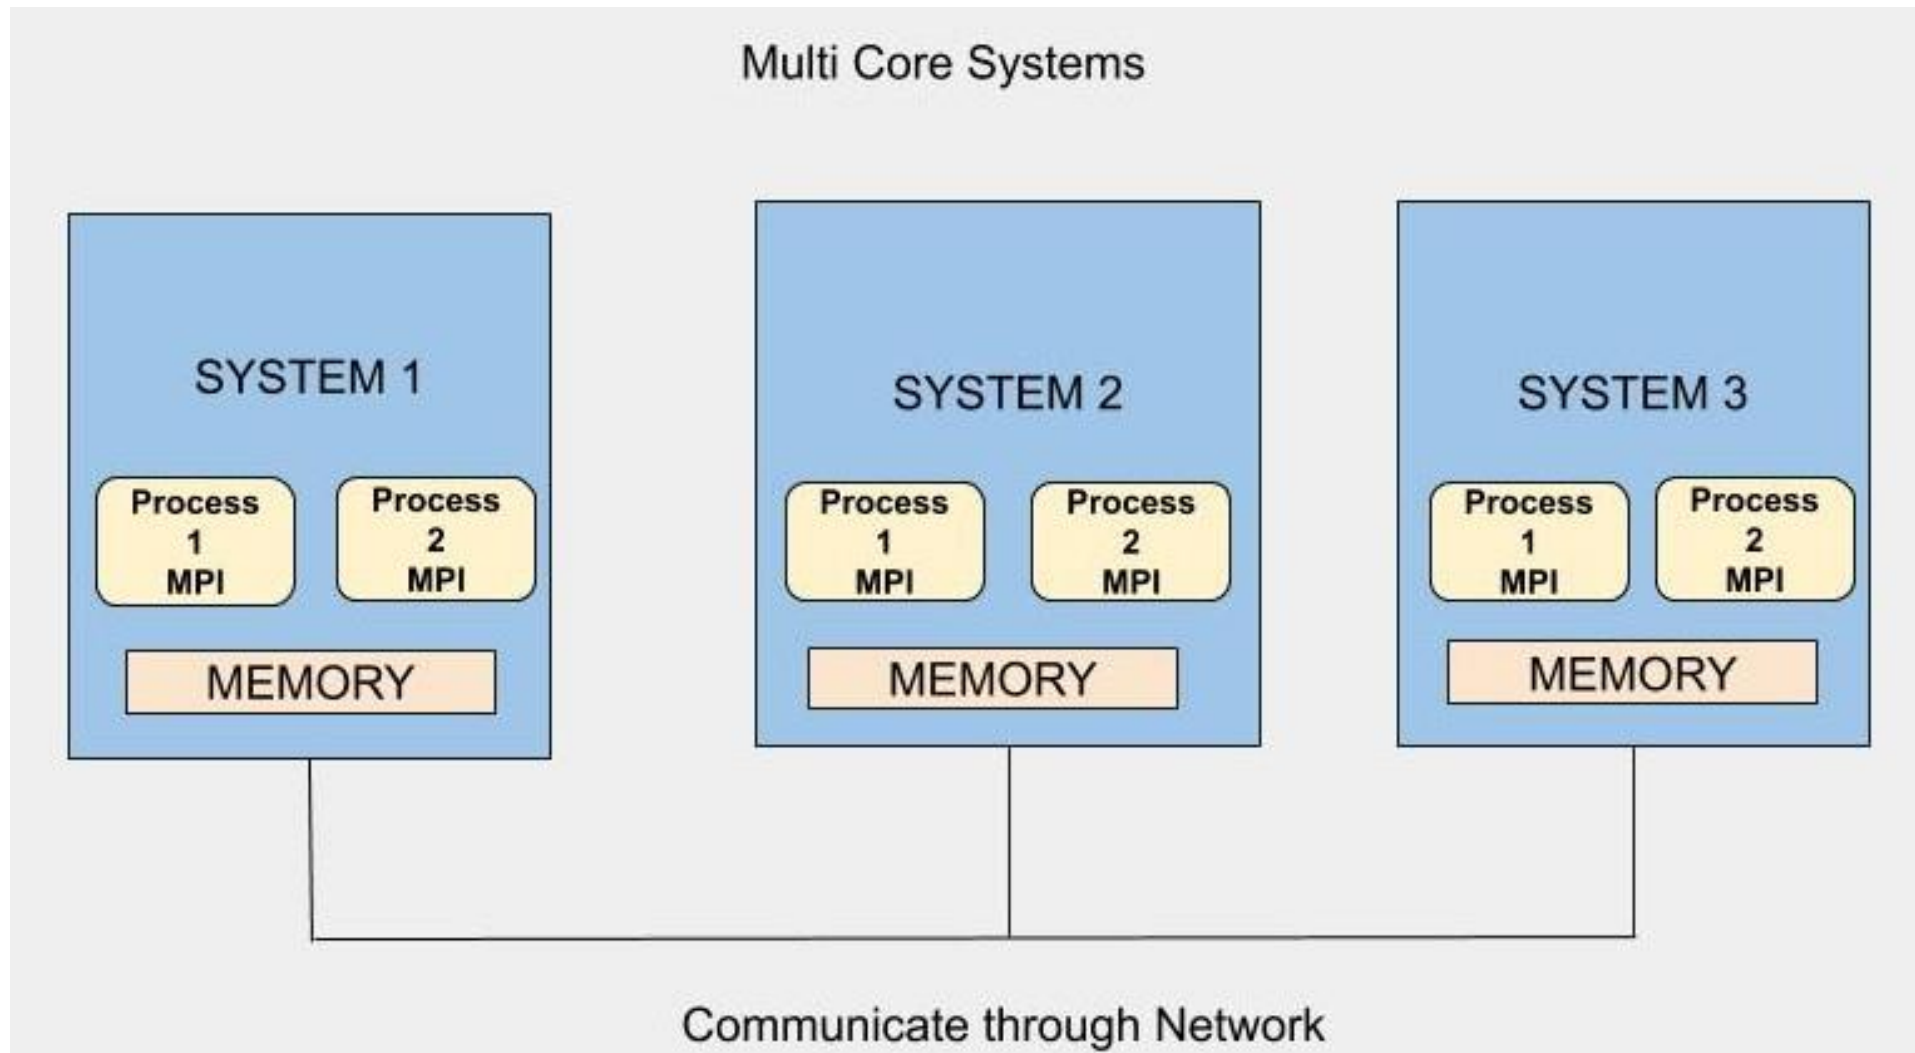
MPI : Major Goals

- Portability :
 - An MPI library exists on ALL parallel computing platforms so it is highly portable.
- Support heterogeneity
- High performance through efficient implementations
- Encourage overlap of communication and computations.
- Reliability

MPI is a Middleware



MPI is a Middleware



MPI Implementations

- OpenMPI (www.open-mpi.org)
- MPICH (www.mpich.org)
- HP MPI
- Intel MPI
- Scali MPI
- IBM MPI

Outline

- Distributed Memory Architecture
- Introduction to MPI
- **Structure of MPI program**
- **Types of Message Passing**
- Basic Routines in Point to Point Communication
- Example programs on Point to Point Communication
- Basic Routines in Collective Communication
- Sample Programs on Collective Communication

STRUCTURE OF MPI PROGRAM

MPI Include File

Initialize MPI Environment

Computations and Message Passing

Terminate MPI Environment

MPI Routines

- **Start and terminate :**
 - To initialize and terminate the MPI environment
- **Communicators :**
 - To identify the communication world (cluster of processes)
- **Getting Information :**
 - To get the number of processes and process ids
- **Sending and Receiving messages :**
 - Actual computation and communication

STRUCTURE OF MPI PROGRAM

MPI Include File

```
#include<mpi.h>
```

Initialize MPI Environment

```
MPI_Init(&argc,&argv);
```

Computations and Message Passing

Terminate MPI Environment

```
MPI_Finalize();
```


MPI Start and Terminate Routines

```
#include<stdio.h>
int main(int argc,char **argv)
{
    -----
    -----
    MPI_Init(&argc,&argv);
    -----
    -----
    MPI_Finalize();
    -----
    return 0;
}
```

Communicators

- MPI defines **communication domain** – set of processes that can communicate with each other.
- **MPI_comm** : data type – stores information about communication domains.
- Default communicator - **MPI_COMM_WORLD**

Communication
Domain



Getting Information

- MPI_Comm_size
- MPI_Comm_rank
- **Syntax :**
- **int MPI_Comm_size(MPI_Comm comm, int *size)**
- **int MPI_Comm_rank(MPI_Comm comm, int *rank)**

General MPI Program

```
#include<mpi.h>
```

```
int main(int argc,char **argv)
```

```
{
```

```
-----
```

```
-----
```

```
MPI_Init(&argc,&argv);
```

```
-----
```

```
MPI_Comm_size(MPI_COMM_WORLD,&size);
```

```
MPI_Comm_rank(MPI_COMM_WORLD,&rank);
```

```
-----
```

```
MPI_Finalize();
```

```
-----
```

```
return 0;
```

```
}
```

Example: Hello World

```
#include<mpi.h>
```

```
int main(int argc,char *argv[ ])
```

```
{
```

```
int size,myrank;
```

```
MPI_Init(&argc,&argv);
```

```
MPI_Comm_size(MPI_COMM_WORLD,&size);
```

```
MPI_Comm_rank(MPI_COMM_WORLD,&myrank);
```

```
printf("Process %d of %d, Hello World",myrank,size);
```

```
MPI_Finalize();
```

```
return 0;
```

```
}
```

MPI Hello World :

```
tans@tans-Inspiron-3542:~/PC$ mpiexec -n 5 ./a.out
Process 0 of 5, Hello World
Process 1 of 5, Hello World
Process 4 of 5, Hello World
Process 2 of 5, Hello World
Process 3 of 5, Hello World
```

MPI Include File



Initialize MPI Environment



Computations and Message Passing



Terminate MPI Environment



Types of Message Passing:

- **Point to Point**

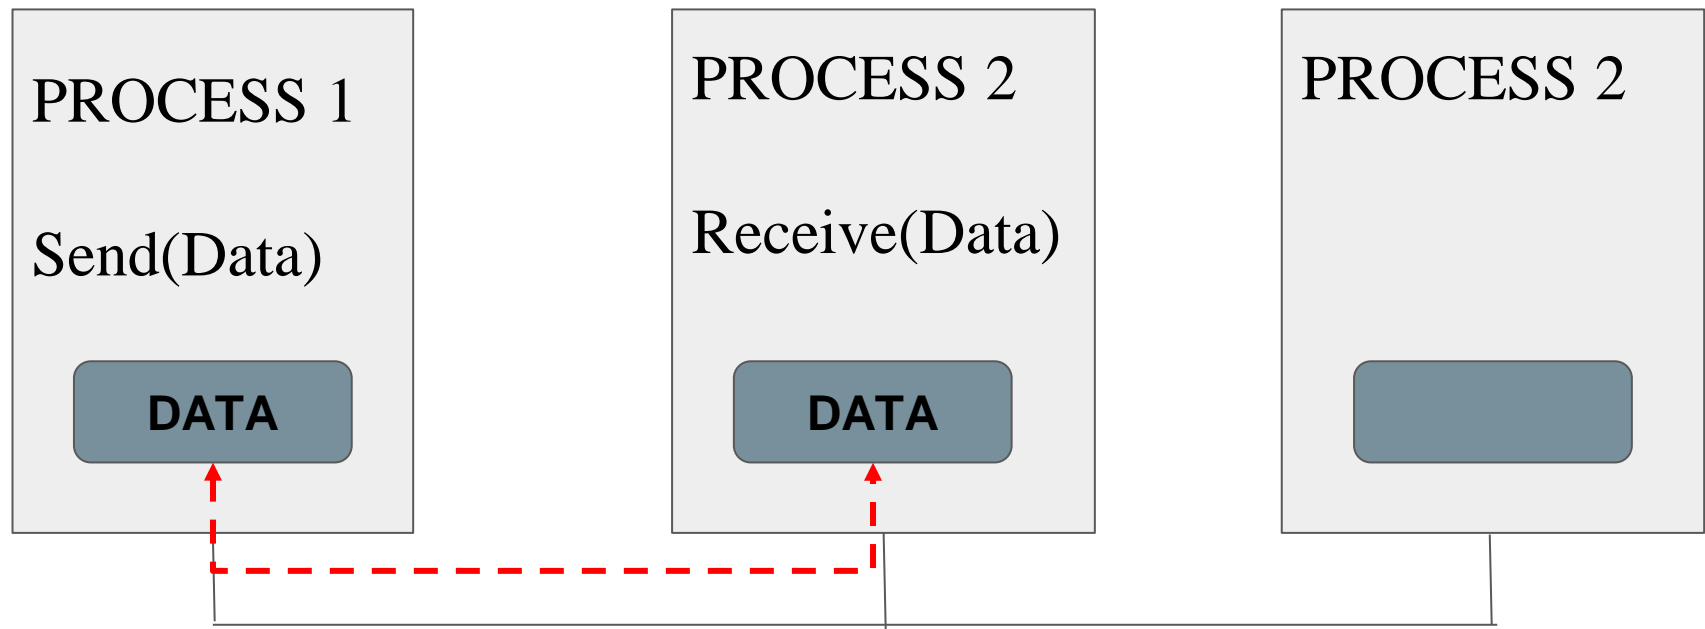
- Two processes
- Send and Receive are the basic functions

- **Collective messages**

- Group of processes involved in communication
- Functions like Broadcast, Scatter, Gather, Parallel Reduction

Point to Point Communication

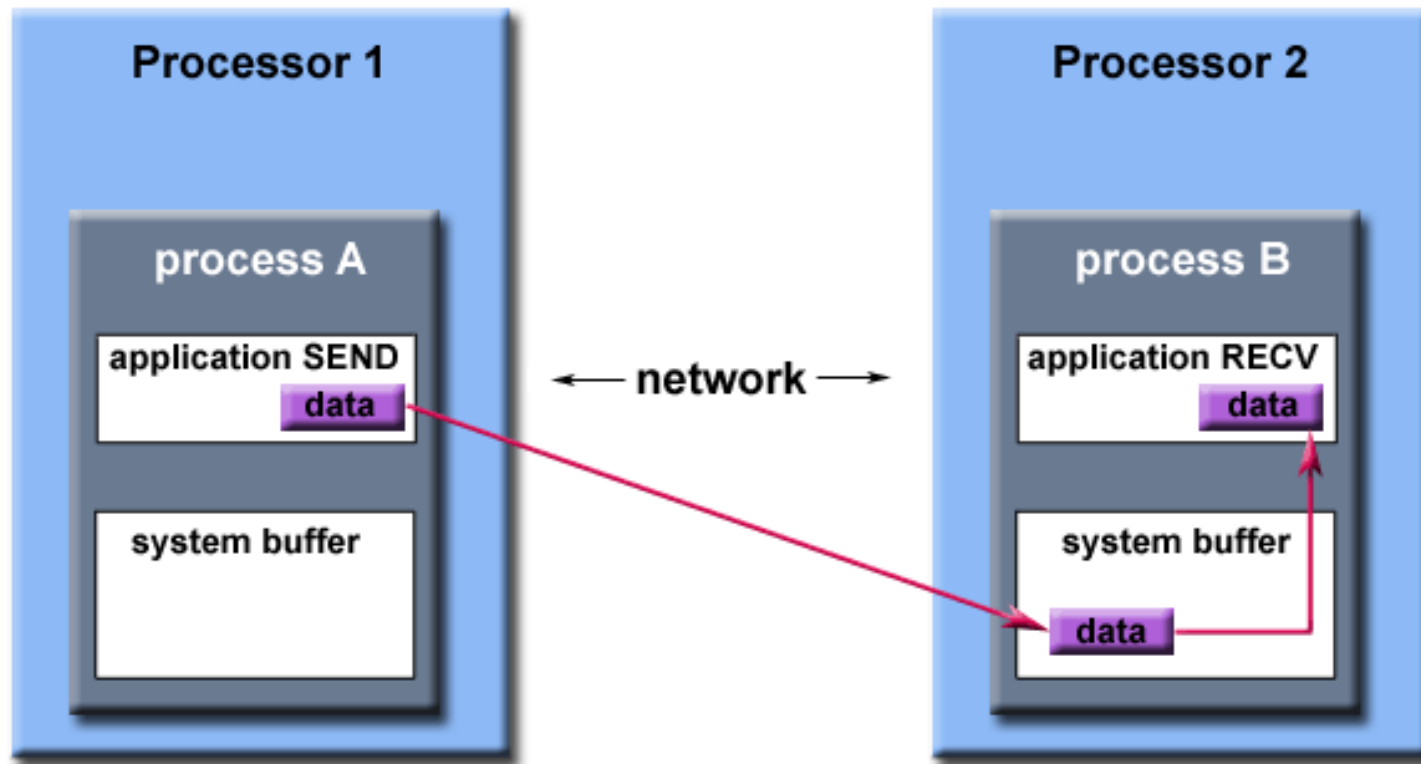
- Two processes involved in sending and receiving data.



- ID of sender and receiver is required.
- Specify what has to be sent and received.
- Communication needs to be synchronized.
- Communication makes use of buffers.

Point to Point Communication

- Data transfer from Sender Process to Receiver Process.



Path of a message buffered at the receiving process

Send and Receive Variants

- **Blocking Send and Receive**
- **Non Blocking Send and Receive**
- **Based on modes of Communication:**
 - Standard
 - Synchronous
 - Buffered
 - Ready

Blocking Send and Receive

- **Basic Send and Receive routine for point to point communication.**
- **MPI Routines:**
 - MPI_Send()
 - MPI_Recv()

Outline

- Distributed Memory Architecture
- Introduction to MPI
- Structure of MPI program
- Types of Message Passing
- **Basic Routines in Point to Point Communication**
- **Example programs on Point to Point Communication**
- **Basic Routines in Collective Communication**
- **Sample Programs on Collective Communication**

Blocking Send and Receive

- **MPI_Send()**

MPI_Send (void *buf, int count, MPI_Datatype type, int dest, int tag, MPI_Comm comm)

Parameters:

- buf** : initial address of send buffer
- count** : number of elements in send buffer (nonnegative integer)
- datatype** : datatype of each send buffer element. Ex : MPI_INT, MPI_CHAR
- dest** : rank of destination (integer)
- tag** : message tag (integer). For tagging send and receive.
- comm** : Communication domain of the communicating processes.

Blocking Send and Receive

- **MPI_Recv():**

MPI_Recv(void ***buf**, int **count**, **MPI_Datatype** **datatype**, int **source**, int **tag**, **MPI_Comm** **comm**, **MPI_Status** ***status**)

Parameters:

- buf** : initial address of receive buffer
- count** : max number of elements in receive buffer (nonnegative integer)
- datatype** : datatype of each receive buffer element. Ex : MPI_INT, MPI_CHAR
- source** : rank of source (integer)
- tag** : message tag (integer). For tagging send and receive.
- comm** : Communication domain of the communicating processes.
- status** : status object (Status). It is a structure containing information about **source**, **tag** and **error code**.

- **MPI DATATYPES:**

Table 1: Basic C datatypes in MPI

MPI Datatype	C data type
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	
MPI_PACKED	

General MPI Program

```
#include<mpi.h>  
int main(int argc,char **argv)  
{  
    ...  
    MPI_Init(&argc,&argv);  
    ...  
    MPI_Comm_size(MPI_COMM_WORLD,&size);  
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);
```

COMPUTATIONS AND MESSAGE PASSING

```
    MPI_Finalize();  
    ...  
    return 0;  
}
```

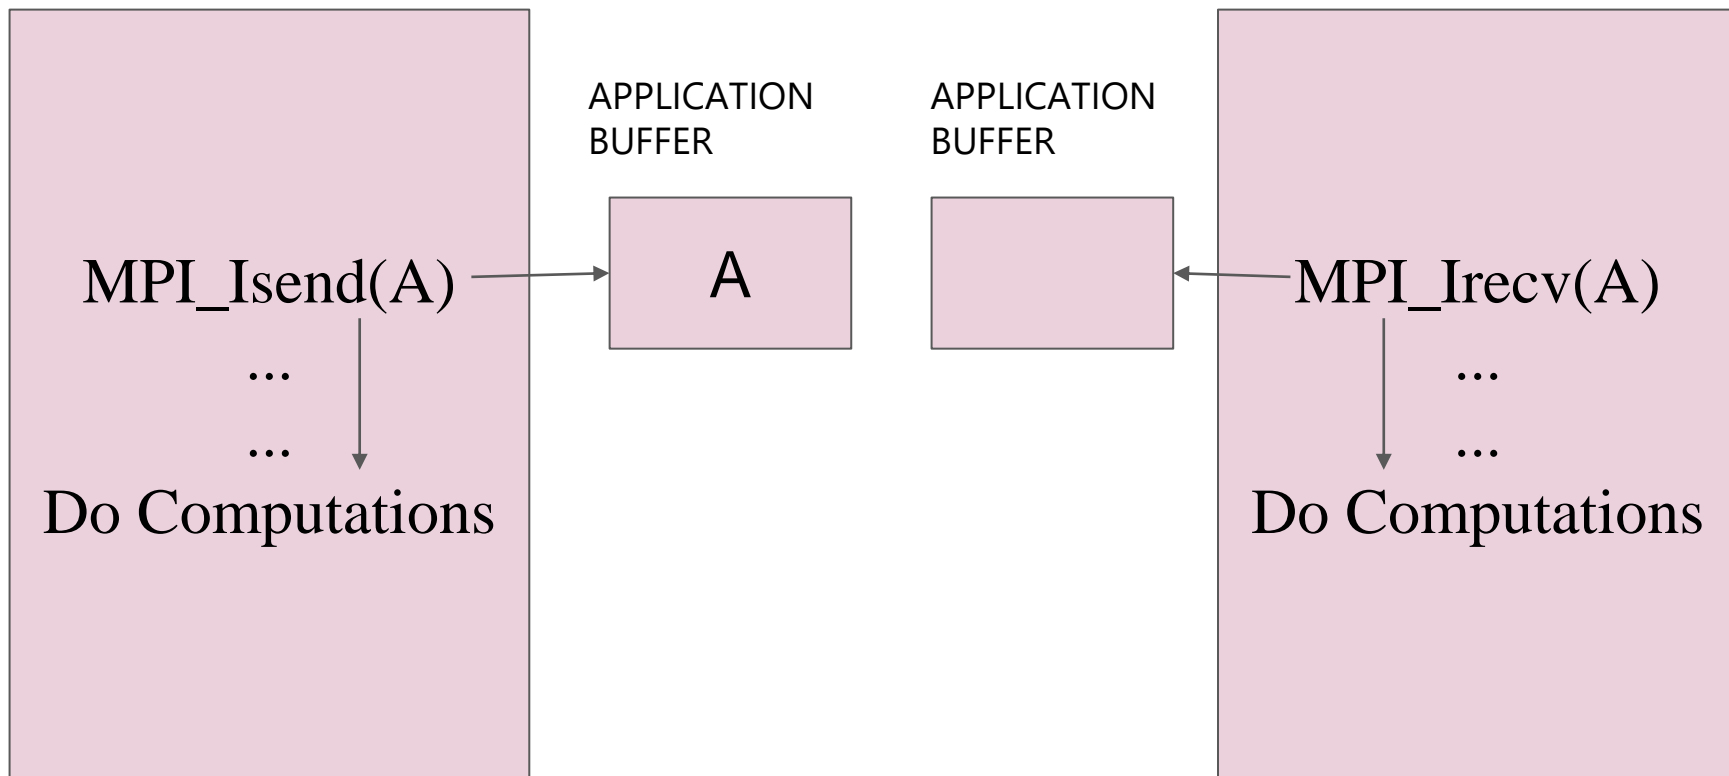
MPI Example - 1

```
for(i=0;i<50;i++) // Process 0 initializes array x
x[i]=i+1;
if(myrank==0)
MPI_Send(x,10,MPI_INT,1,1,MPI_COMM_WORLD);
else if(myrank==1)
{
MPI_Recv(y,10,MPI_INT,0,1,MPI_COMM_WORLD,&status);
printf("Process %d Received Data from Process %d\n",
myrank,status.MPI_SOURCE);
for(i=0;i<10;i++)
printf("%d\t",y[i]);
}
Process 1 Recieved data from Process 0
```

1	2	3	4	5	6	7	8	9	10
---	---	---	---	---	---	---	---	---	----

Non Blocking Send and Receive

- Allows overlapping of computation and communication
- Advantage is Performance Gain



Non Blocking Send and Receive

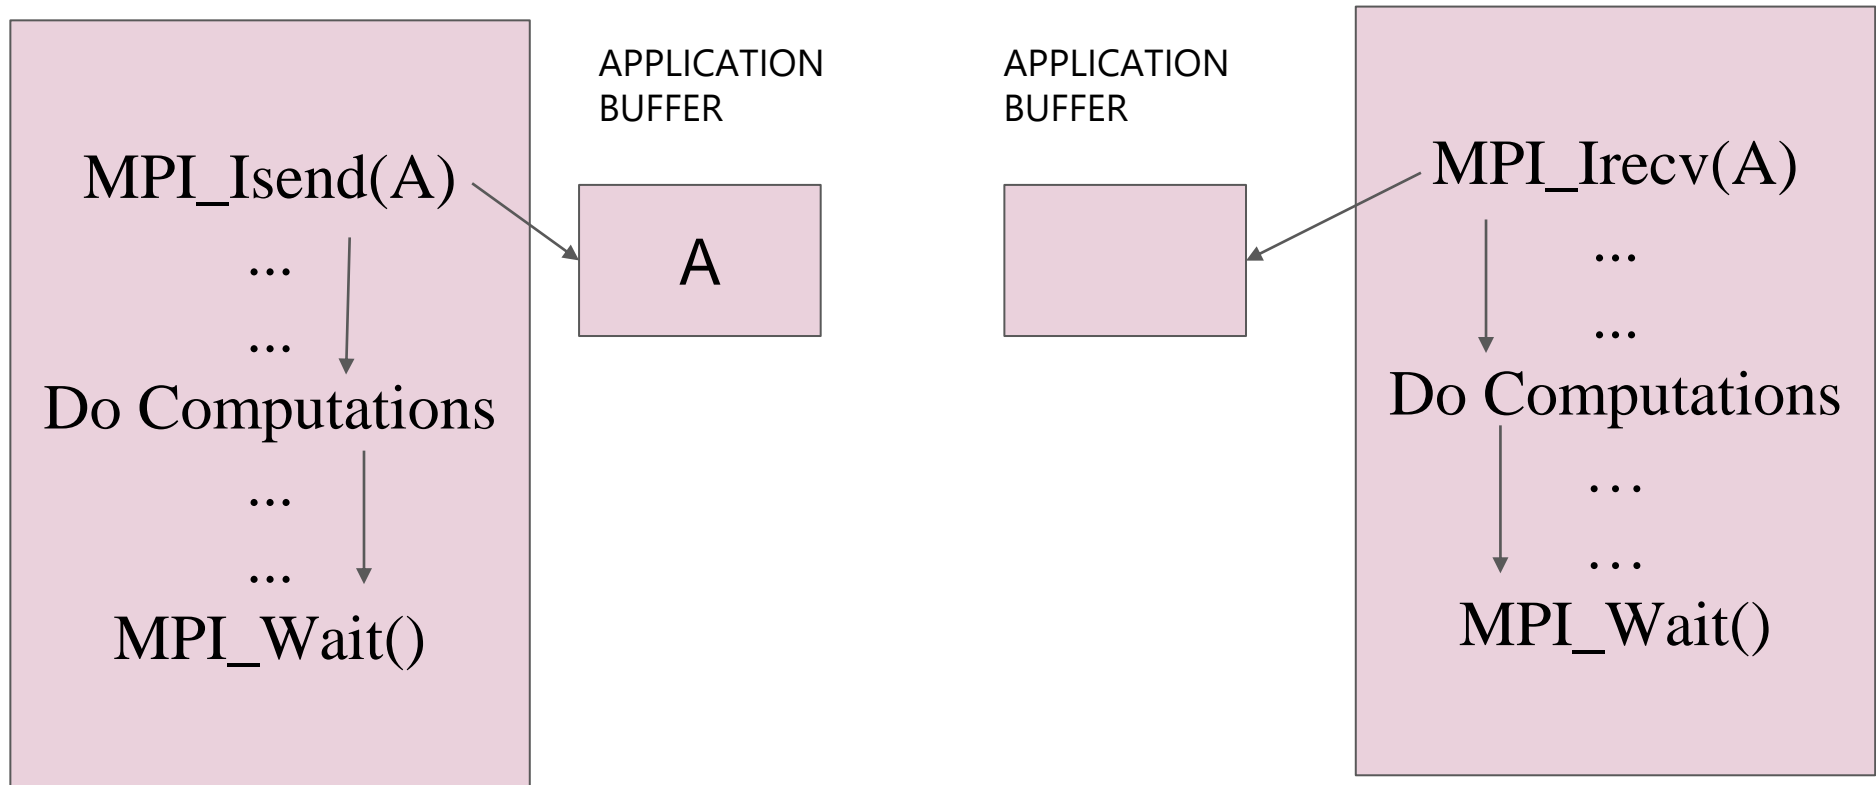
MPI_Isend (&buf,count,datatype,dest,tag,comm,&request)

MPI_Irecv (&buf,count,datatype,source,tag,comm,&request)

Parameters:

- Same as Send() and Recv() except for request
- **request** : handle. This helps to get information about MPI_Isend and MPI_Irecv status.
- Used in routines : **MPI_Wait()** and **MPI_Test()**

Non Blocking Send and Receive



MPI_Wait() and MPI_Test()

Syntax :

```
int MPI_Wait( MPI_Request *request, MPI_Status *status );
```

```
int MPI_Test( MPI_Request *request, int *flag, MPI_Status *status );
```

- If request is set to MPI_REQUEST_NULL (set if operation is completed) then:
 - MPI_Wait returns immediately with an empty status.
 - MPI_Test sets flag to true and returns an empty status.

MPI Example - 2

```
if(myrank==0)
{
x=10;
MPI_Isend(&x,1,MPI_INT,1,20,MPI_COMM_WORLD,&request);
printf("Send returned immediately\n");
}
else if(myrank==1)
{
MPI_Irecv(&x,1,MPI_INT,0,25,MPI_COMM_WORLD,&request);
printf("Receive returned immediately\n");
printf("Process %d of %d, Value of x is %d\n",myrank,size,x);
}
}
```

```
tans@tans-Inspiron-3542:~/PC$ mpiexec -n 2 ./a.out
Send returned immediately
Receive returned immediately
Process 1 of 2, Value of x is 0
```


What is the risk here?

```
if(myrank==0)
{
    x=10;

    MPI_Isend(&x,1,MPI_INT,1,20,MPI_COMM_WORLD,&request);

    printf("Send returned immediately\n");

    x=x+10;

}
```

Make sure that x is available for reuse:

```
if(myrank==0)
```

```
{
```

```
x=10;
```

```
MPI_Isend(&x,1,MPI_INT,1,20,MPI_COMM_WORLD,&request);
```

```
printf("Send returned immediately\n");
```

```
MPI_Wait(request, status)
```


```
x=x+10;
```

```
}
```

Communication Modes

- **Standard Mode** : Calls block until message has been either transferred or copied to an internal buffer for later delivery. Ex: `MPI_Send()` and `MPI_Recv()`
- **Buffered Mode** : Send may start and return before a matching receive. `MPI_Bsend()`
- **Synchronous Mode** : Call blocks until matching receive has been posted and the message reception has started. `MPI_Ssend()`
- **Ready Mode** : Requires that a matching receive is already posted. `MPI_Rsend()`.

Buffered Mode




MPI_BUFFER_ATTACH(buffer, size)	
buffer	initial buffer address (choice)
size	buffer size, in bytes (integer)

NOTE: A user may specify a buffer to be used for buffering messages sent in buffered mode.


Synchronous Mode

We see that the data is not copied to system buffer.



Ready Mode

We make use of `MPI_Barrier()` to wait for the receive to be posted. This will not result in error.



MPI-Example - 3

```
if(myrank==0) {
```

```
// Blocking send will expect matching receive at the destination In Standard mode, Send will  
return after copying the data to the buffer
```

```
MPI_Send(x,10,MPI_INT,1,1,MPI_COMM_WORLD);
```

```
// This send will be initiated and matching receive is already there so the program will not  
lead to deadlock
```

```
MPI_Send(y,10,MPI_INT,1,2,MPI_COMM_WORLD);
```

```
}
```

```
else if(myrank==1)
```

```
{
```

```
// P1 will block as it has not received a matching send with tag 2
```

```
MPI_Recv(x,10,MPI_INT,0,2,MPI_COMM_WORLD,&status);
```

```
MPI_Recv(y,10,MPI_INT,0,1,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
```

```
}
```

MPI Example 3

PROCESS 1

MPI_Send(x,10,..1,1,..);

MPI_Send(y,10,..,1,2,..);

PROCESS 2

MPI_Recv(x,10,..,0,2,..,..);

 BLOCK

MPI_Recv(y,10,..,0,1,..,..);

MPI Example - 4

```
if(myrank==0) {
```

```
    MPI_Ssend(x,10,MPI_INT,1,1,MPI_COMM_WORLD);
```

```
    MPI_Send(y,10,MPI_INT,1,2,MPI_COMM_WORLD);
```

```
}
```

```
else if(myrank==1)
```

```
{
```

```
    MPI_Recv(x,10,MPI_INT,0,2,MPI_COMM_WORLD,&status);
```

```
MPI_Recv(y,10,MPI_INT,0,1,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
```

```
}
```

MPI Example - 4

```
if(myrank==0) {
```

```
MPI_Ssend(x,10,MPI_INT,1,1,MPI_COMM_WORLD);
```

// Synchronous Blocking send will expect matching receive at the destination.
This results in deadlock.

```
    MPI_Send(y,10,MPI_INT,1,2,MPI_COMM_WORLD); // This call will not be  
    executed  
}
```

```
else if(myrank==1)
```

```
{
```

```
    MPI_Recv(x,10,MPI_INT,0,2,MPI_COMM_WORLD,&status); // P1 will block  
    as it has not received a matching send with tag 2
```

```
MPI_Recv(y,10,MPI_INT,0,1,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
```

```
}
```


Outline

- Distributed Memory Architecture
- Introduction to MPI
- Structure of MPI program
- Types of Message Passing
- **Basic Routines in Point to Point Communication**
- **Example programs on Point to Point Communication**
- **Basic Routines in Collective Communication**
- **Sample Programs on Collective Communication**

Communication Modes

- **Standard Mode** : Calls block until message has been either transferred or copied to an internal buffer for later delivery. Ex: `MPI_Send()` and `MPI_Recv()`
- **Buffered Mode** : Send may start and return before a matching receive. `MPI_Bsend()`
- **Synchronous Mode** : Call blocks until matching receive has been posted and the message reception has started. `MPI_Ssend()`
- **Ready Mode** : Requires that a matching receive is already posted. `MPI_Rsend()`.

Buffered Mode




MPI_BUFFER_ATTACH(buffer, size)	
buffer	initial buffer address (choice)
size	buffer size, in bytes (integer)

NOTE: A user may specify a buffer to be used for buffering messages sent in buffered mode.


Synchronous Mode

We see that the data is not copied to system buffer.



Ready Mode

We make use of `MPI_Barrier()` to wait for the receive to be posted. This will not result in error.




Outline

- Distributed Memory Architecture
- Introduction to MPI
- Structure of MPI program
- Types of Message Passing
- Basic Routines in Point to Point Communication
- Example programs on Point to Point Communication
- **Basic Routines in Collective Communication**
- **Sample Programs on Collective Communication**

Collective Communication

- Multiple processes in same communicator involve in collective communication.
- They are blocking calls.
- No tags required.



Collective Communication

- **Barrier**
- **Broadcast**
- **Scatter**
- **Gather**
- **Reduce**
- **Scatterv**
- **Gatherv**


Collective communication: MPI_Barrier

- Mainly used for synchronization
- The call returns only after all the processes have called Barrier function.
- Uses:
 - Access to files
 - Achieve consistency

Syntax: MPI_Barrier(MPI_COMM_WORLD)

Collective Communication: Broadcast

- **MPI_Bcast(buf, count, datatype, source, comm)**
 - buf : send buffer of sender and receive buffer of receiver
 - source : process which sends data to others

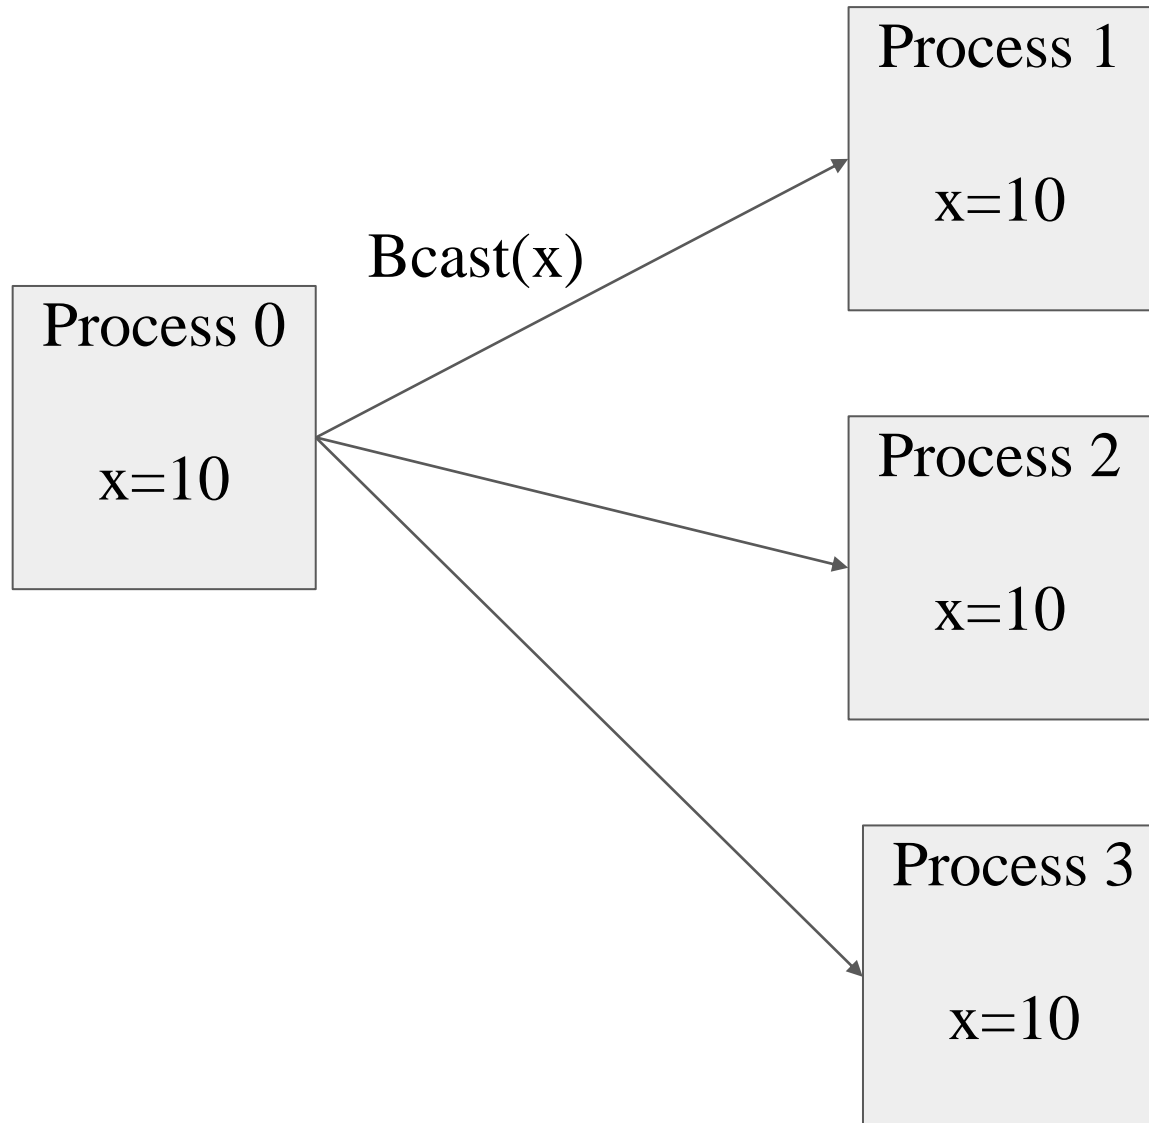


broadcast

MPI Example - 5

```
if(myrank==0)
{
scanf("%d",&x);
}
MPI_Bcast(&x,1,MPI_INT,0,MPI_COMM_WORLD);
printf("Value of x in process %d : %d\n",myrank,x);
MPI_Finalize();
return 0;
}
```


Bcast():



Broadcast Output:

```
tans@tans-Inspiron-3542:~/PC$ mpiexec -n 4 ./a.out
3
Value of x in process 0 : 3
Value of x in process 1 : 3
Value of x in process 2 : 3
Value of x in process 3 : 3
```

Collective Communication: Scatter



Collective Communication: Scatter

MPI_Scatter(sendbuf, sendcount, datatype, recvbuf, recvcount, datatype, root, comm)

Parameters:

sendbuf : sender buffer


sendcount : specify the number of elements to be sent. **recvcount** should be same as **sendcount**

recvbuf : recv buffer

root : Sender

MPI_Scatter

Example:



MPI Example - 6


```
if(myrank==0)
{
printf("Enter values into array x:\n");
for(i=0;i<8;i++)
scanf("%d",&x[i]);
}
MPI_Scatter(x,2,MPI_INT,y,2,MPI_INT,0,MPI_COMM_WORLD);
for(i=0;i<2;i++)
printf("\nValue of y in process %d : %d\n",myrank,y[i]);
```

Output

```
tans@tans-Inspiron-3542:~/PC$ mpiexec -n 4 ./a.out
Enter values into array x:
1 2 3 4 5 6 7 8

Value of y in process 0 : 1
Value of y in process 0 : 2
Value of y in process 1 : 3
Value of y in process 1 : 4
Value of y in process 2 : 5
Value of y in process 2 : 6
Value of y in process 3 : 7
Value of y in process 3 : 8
```

Collective Communication: Gather



Collective Communication: Gather

MPI_Gather(sendbuf, sendcount, datatype, recvbuf, recvcount, datatype, root, comm)

Parameters:


sendbuf: buffer of sending processes

sendcount and recvcount value is same

recvbuf: root process's buffer

root : process where the data is gathered

MPI_Gather



MPI-Example 7

x=10, y[50]

```
MPI_Gather(&x,1,MPI_INT,y,1,MPI_INT,0,MPI_COMM_WORLD);
```

// Value of x at each process is copied to array y in Process 0

```
if(myrank==0)
```

```
{
```

```
for(i=0;i<size;i++)
```

```
printf("\nValue of y[%d] in process %d : %d\n",i,myrank,y[i]);
```

```
}
```


Output

```
tans@tans-Inspiron-3542:~/PC$ mpiexec -n 4 ./a.out
```

```
Value of y[0] in process 0 : 10
```

```
Value of y[1] in process 0 : 10
```

```
Value of y[2] in process 0 : 10
```

```
Value of y[3] in process 0 : 10
```

```
tans@tans-Inspiron-3542:~/PC$ mpiexec -n 6 ./a.out
```

```
Value of y[0] in process 0 : 10
```

```
Value of y[1] in process 0 : 10
```

```
Value of y[2] in process 0 : 10
```

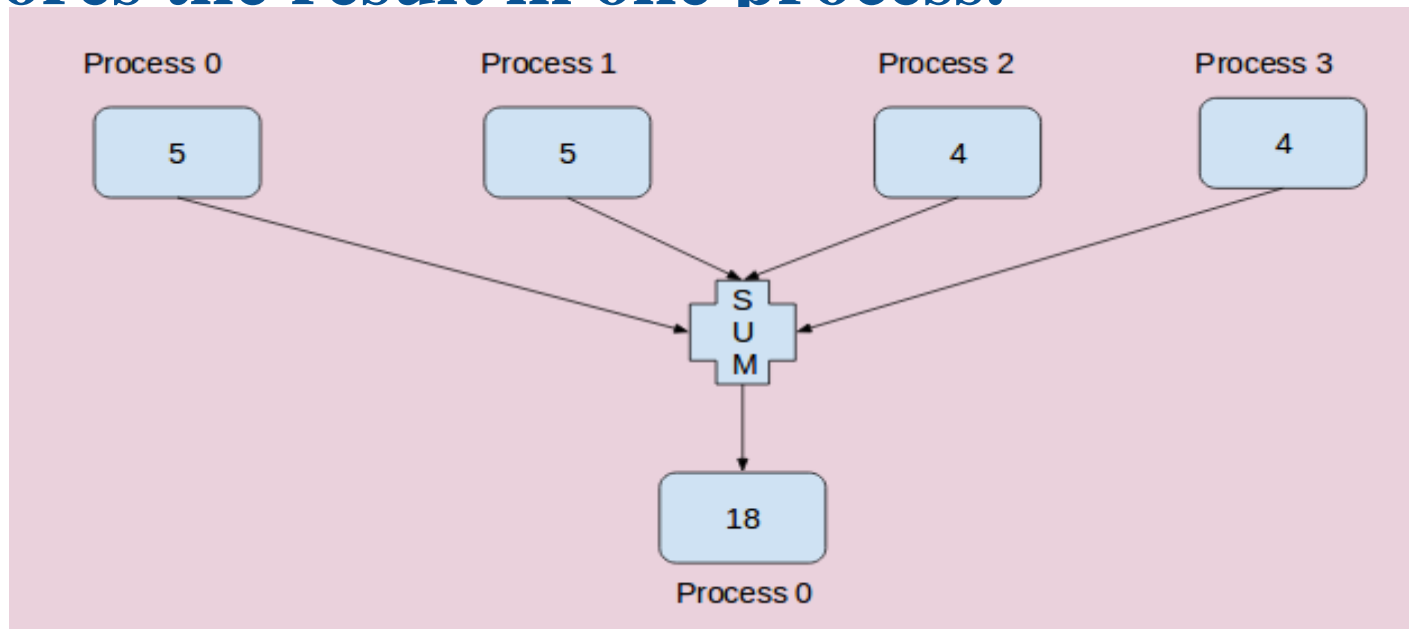
```
Value of y[3] in process 0 : 10
```

```
Value of y[4] in process 0 : 10
```

```
Value of y[5] in process 0 : 10
```

Collective Communication: Reduce

- Allows to perform computations on data present at multiple processes.
- Computations like : Sum, Product, Maximum, Minimum
- Stores the result in one process.



Collective Communication: Reduce

MPI_Reduce(sendbuf, recvbuf, count, datatype, operation, dest, comm)

Parameters:

count: size of receive buffer

operation:

MPI name	Operation
MPI_MAX	Maximum
MPI_MIN	Minimum
MPI_SUM	Summation
MPI_PROD	Product
MPI LAND	Logical AND
MPI_LOR	Logical OR
MPI_LXOR	Logical XOR

MPI Example - 8

```
x=myrank;  
MPI_Reduce(&x,&y,1,MPI_INT,MPI_SUM,0,MPI_COMM_WORLD)  
;  
if(myrank==0)  
{  
printf("Value of y after reduce : %d\n",y);  
}
```

Output

```
tans@tans-Inspiron-3542:~/PC$ mpiexec -n 3 ./a.out  
Value of y after reduce : 3  
tans@tans-Inspiron-3542:~/PC$ mpiexec -n 4 ./a.out  
Value of y after reduce : 6
```


Outline

- Distributed Memory Architecture
- Introduction to MPI
- Structure of MPI program
- Types of Message Passing
- Basic Routines in Point to Point Communication
- Example programs on Point to Point Communication
- **Basic Routines in Collective Communication**
- **Sample Programs on Collective Communication**


MORE Collective Communication Routines:

- **MPI_Gatherv()**
- **MPI_Scatterv()**
- **MPI_Allgather**
- **MPI_AllReduce()**
- **MPI_Scan()**
- **MPI_Comm_Split()**

MPI_Allgather



MPI_Scatterv



MPI_Scatterv():

MPI_Scatterv(sendbuf, sendcounts, displacement, datatype, recvbuf, recvcount, datatype, root, comm)

Parameters:

sendcounts : array with number of elements to be sent to each process. ex: sendcount[0]=10 means send 10 elements to Process zero. sendcount[1]=20 means send 20 elements to Process one.

displacement: array which holds the index from where the data is to be sent to each process. Ex: disp[0]=0 means Process zero gets elements starting with index zero. disp[1]=10 means Process 1 will get elements starting from index 10.

MPI_Scatterv

```
9 if(myrank==0)
10 {
11     printf("Enter the number of elements:\n");
12     scanf("%d",&m);
13     printf("Initializing array x:\n");
14     for(i=0;i<m;i++)
15         x[i]=i+1;
16     disp[0]=0;
17     for(i=0;i<size;i++)
18     {
19         z[i]=i+2; // Process 0 gets 2 elements, Process 1 gets 3, Process 2 gets 4 and so on
20         disp[i+1]=disp[i]+z[i]; // disp[1]=2, disp[2]=5, disp[3]=9 and so on
21     }
22 }
23
24 MPI_Scatterv(x,z,disp,MPI_INT,y,myrank+2,MPI_INT,0,MPI_COMM_WORLD);
25
```


OUTPUT:

```
Enter the number of elements:
20

Value of y in process 1 : 3
Value of y in process 1 : 4
Value of y in process 1 : 5
Value of y in process 2 : 6
Value of y in process 2 : 7
Value of y in process 2 : 8
Value of y in process 2 : 9
Value of y in process 3 : 10
Value of y in process 3 : 11
Value of y in process 3 : 12
Value of y in process 3 : 13
Value of y in process 3 : 14
Value of y in process 4 : 15
Value of y in process 4 : 16
Value of y in process 4 : 17
Value of y in process 4 : 18
Value of y in process 4 : 19
Value of y in process 4 : 20
Initializing array x:

Value of y in process 0 : 1
Value of y in process 0 : 2
```

MPI_Gatherv



MPI_Gatherv():

MPI_Gatherv(sendbuf, sendcount, datatype, recvbuf, recvcunts, displacements, datatype, root, comm)

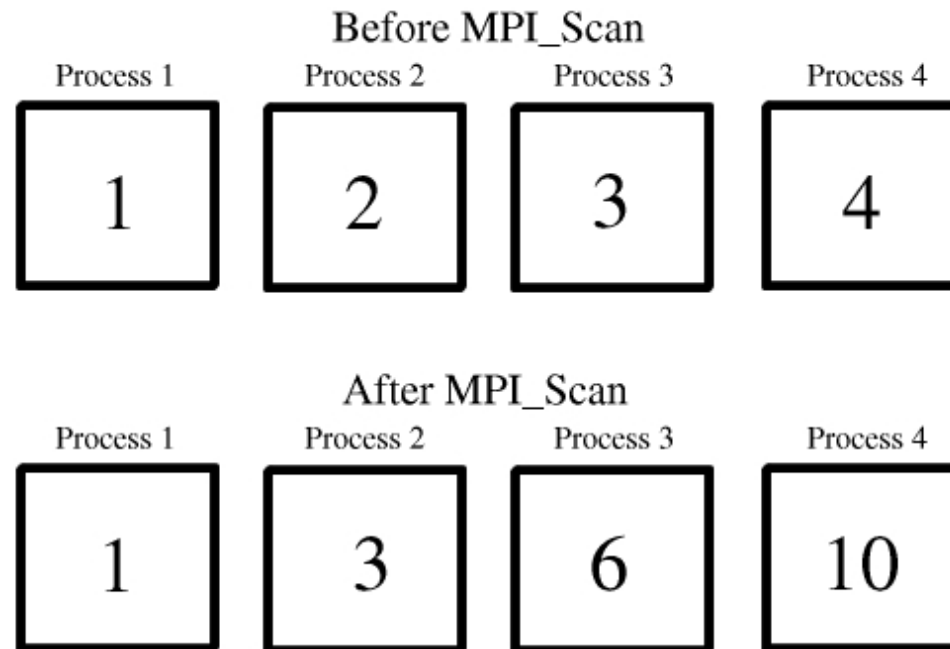
Parameters:

recvcunts : array with number of elements to be received from each process.

displacement: array which holds the beginning index where the data is to be received from each process.

MPI_Scan

`int MPI_Scan(sendbuf, recvbuf, int count, datatype, MPI_Op, comm)`



MPI_Comm_Split : Split the communication Domain

`MPI_Comm_Split(MPI_Comm comm, int color, int key, MPI_Comm *newcomm);`

color: controls subset assignment

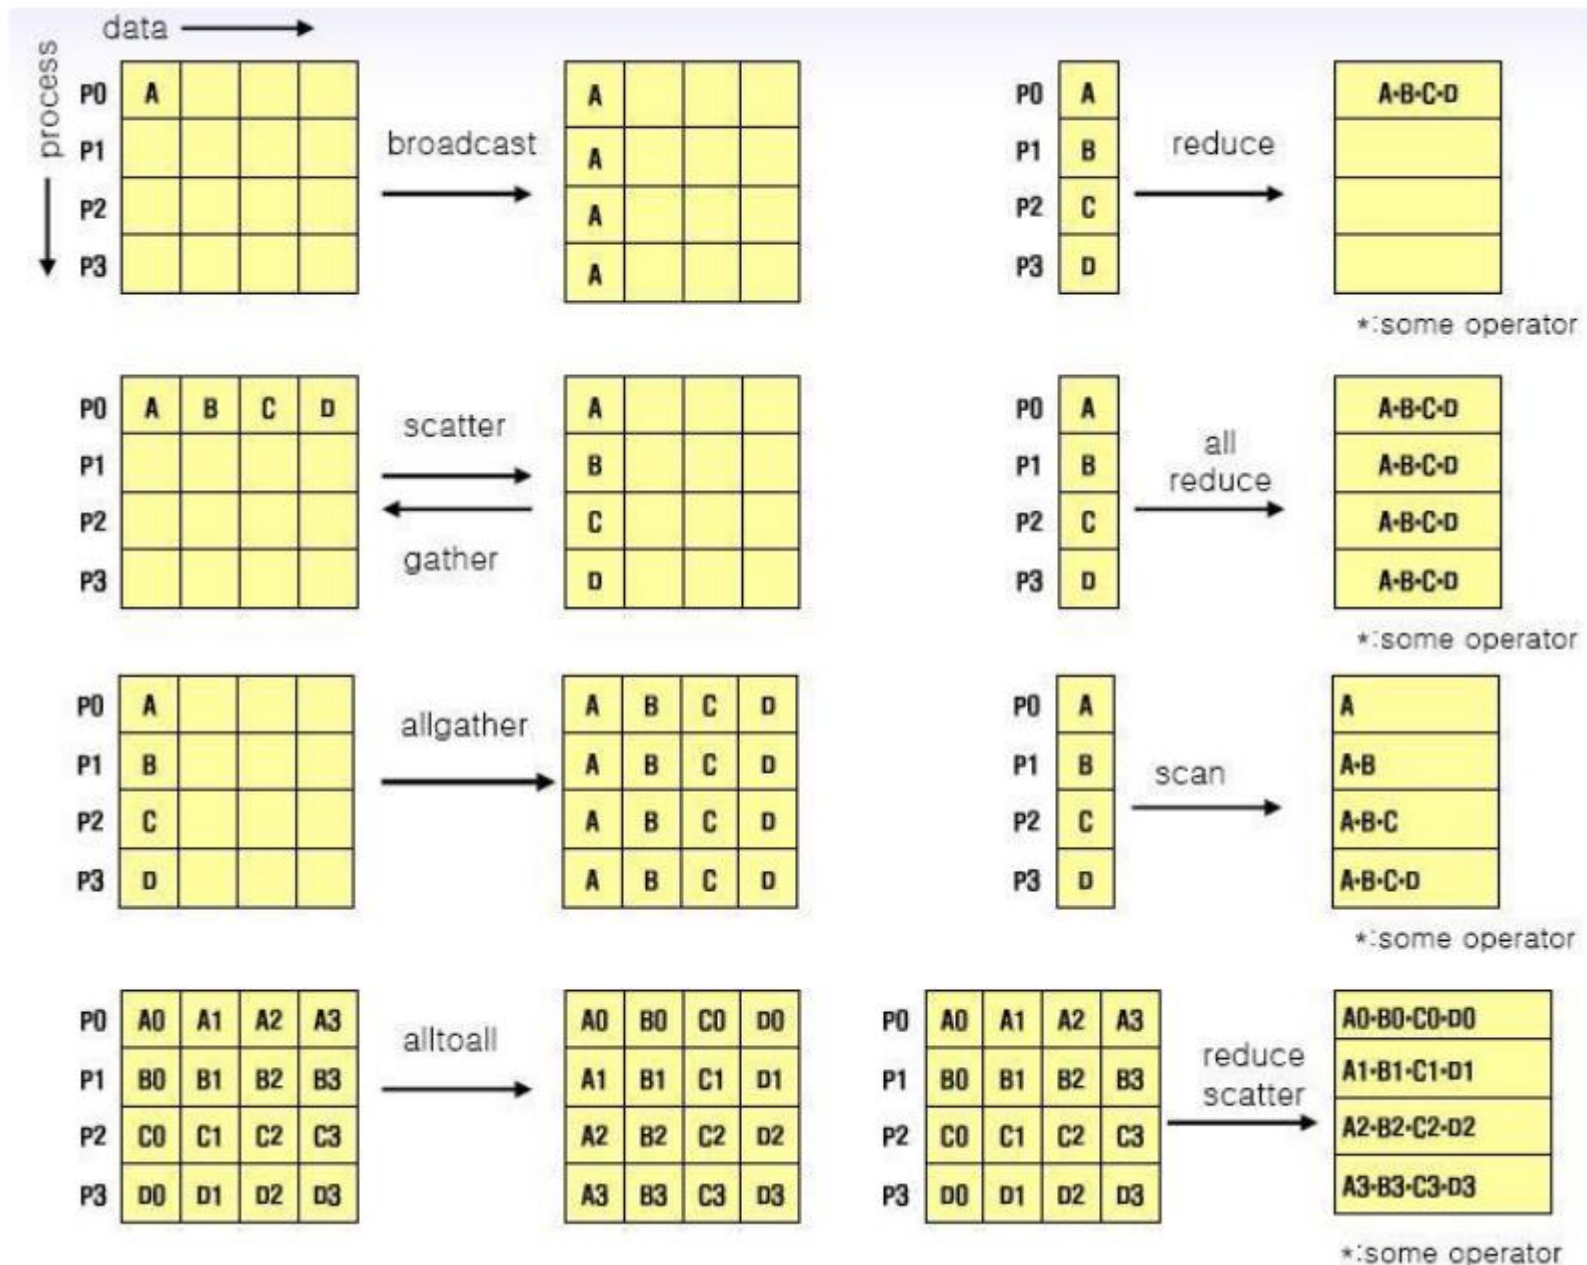
key: controls rank assignment of processes in different group

Ex: `MPI_Comm_split(MPI_COMM_WORLD,0,0,&comm1);`

`MPI_Comm_split(MPI_COMM_WORLD,1,0,&comm2);`

`MPI_Comm_split(MPI_COMM_WORLD,2,0,&comm3);`

MPI Collective Routine



Summary

- MPI provides a simplified way for sending and receiving messages
- MPI rich set of collective functions
- MPI helps for developing Scalable and Portable Parallel Programs
- MPI is the defacto standard for Distributed Memory Parallelism

Thank You