

National Institute of Technology Karnataka Surathkal

Department of Information Technology



IT 301 Parallel Computing

Shared Memory Programming Technique (2)

OpenMP : *parallel, for Clauses*

Dr. Geetha V

Assistant Professor

Dept of Information Technology

NITK Surathkal

Index

- OpenMP

- Program Structure
- Directives : *parallel* , *for*
- Clauses
 - If
 - Private
 - Shared
 - Default(shared/none)
 - Firstprivate
 - Num_threads

- References

Course Outline

Course Plan: Theory:

Part A: Parallel Computer Architectures

Week 1,2,3: ***Introduction to Parallel Computer Architecture:*** Parallel Computing, Parallel architecture, bit level, instruction level , data level and task level parallelism. Instruction level parallelism: pipelining(Data and control instructions), scalar and superscalar processors, vector processors. Parallel computers and computation.

Week 4,5: Memory Models: UMA, NUMA and COMA. Flynn's classification, Cache coherence,

Week 6,7: Amdahl's Law. Performance evaluation, Designing parallel algorithms : Divide and conquer, Load balancing, Pipelining.

Week 8 -11: ***Parallel Programming techniques like Task Parallelism using TBB, TL2, Cilk++ etc. and software transactional memory techniques.***

Course Outline

Part B: OpenMP/MPI/CUDA

Week 1,2,3 : **Shared Memory Programming Techniques:** Introduction to OpenMP :

Directives: *parallel, for, sections, task, single, critical, barrier, taskwait, atomic.*

Clauses: *private, shared, firstprivate, lastprivate, reduction, nowait, ordered, schedule, collapse, num_threads, if().*

Week 4,5: **Distributed Memory programming Techniques:** MPI: Blocking, Non-blocking.

Week 6,7 : CUDA : OpenCL, Execution models, GPU memory, GPU libraries.

Week 10,11,: **Introduction to accelerator programming using CUDA/OpenCL and Xeon-phi. Concepts of Heterogeneous programming techniques.**

Practical:

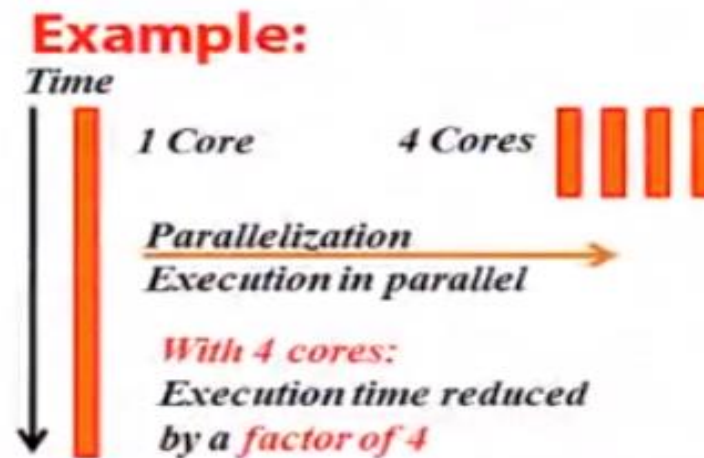
Implementation of parallel programs using OpenMP/MPI/CUDA.

Assignment: Performance evaluation of parallel algorithms (in group of 2 or 3 members)

1. Introduction

- Serial Programming
 - Develop a serial program and Optimize for performance
- Real World scenario:
 - Run multiple programs
 - Large and Complex problems
 - Time consuming
- Solution:
 - Use parallel machines
 - Use Multi-core Machines
- Why Parallel ?
 - Reduce the execution time
 - Run multiple programs

- What is parallel programming?
 - Obtain the same amount of computation with multiple cores or threads at low frequency (Fast)




1. Introduction

- Concurrency

- Condition of a system in which multiple tasks are logically active at the same timebut they may not necessarily run in parallel

- Parallelism

- Subset of concurrency
- Condition of a system in which multiple tasks are active at the same time and run in parallel



1. Introduction


- Shared Memory Machines

- All processors share the same memory
- The variables can be shared or private
- Communication via shared memory

Multi-threading

- Portable, easy to program and use
- Not very scalable

- OpenMP based Programming




- Distributed Memory Machines

- Each processor has its own memory
- The variables are Independent
- Communication by passing messages (network)

Multi-Processing

- Difficult to program
- Scalable

- MPI based Programming



1. OpenMP : API


Open Specification for Multi-Processing (OpenMP)

- Library used to divide computational work in a program and add parallelism to serial program (create threads)
- An Application Program Interface (API) that is used explicitly direct multi-threaded, shared memory parallelism
- **API Components**
 - Compiler Directives
 - Runtime library routines
 - Environment variables
- **Standardization**
 - Jointly defined and endorsed by major computer hardware and software vendors

1. OpenMP

FORK – JOIN Parallelism

- OpenMP program begin as a single process: the master thread. The master thread executes sequentially until the first parallel region construct is encountered.
- When a parallel region is encountered, master thread
 - Create a group of threads by FORK.
 - Becomes the master of this group of threads and is assigned the thread id 0 within the group.
- The statement in the program that are enclosed by the parallel region construct are then executed in parallel among these threads.
- JOIN: When the threads complete executing the statement in the parallel region construct, they synchronize and terminate, leaving only the master thread.



1. OpenMP

- I/O

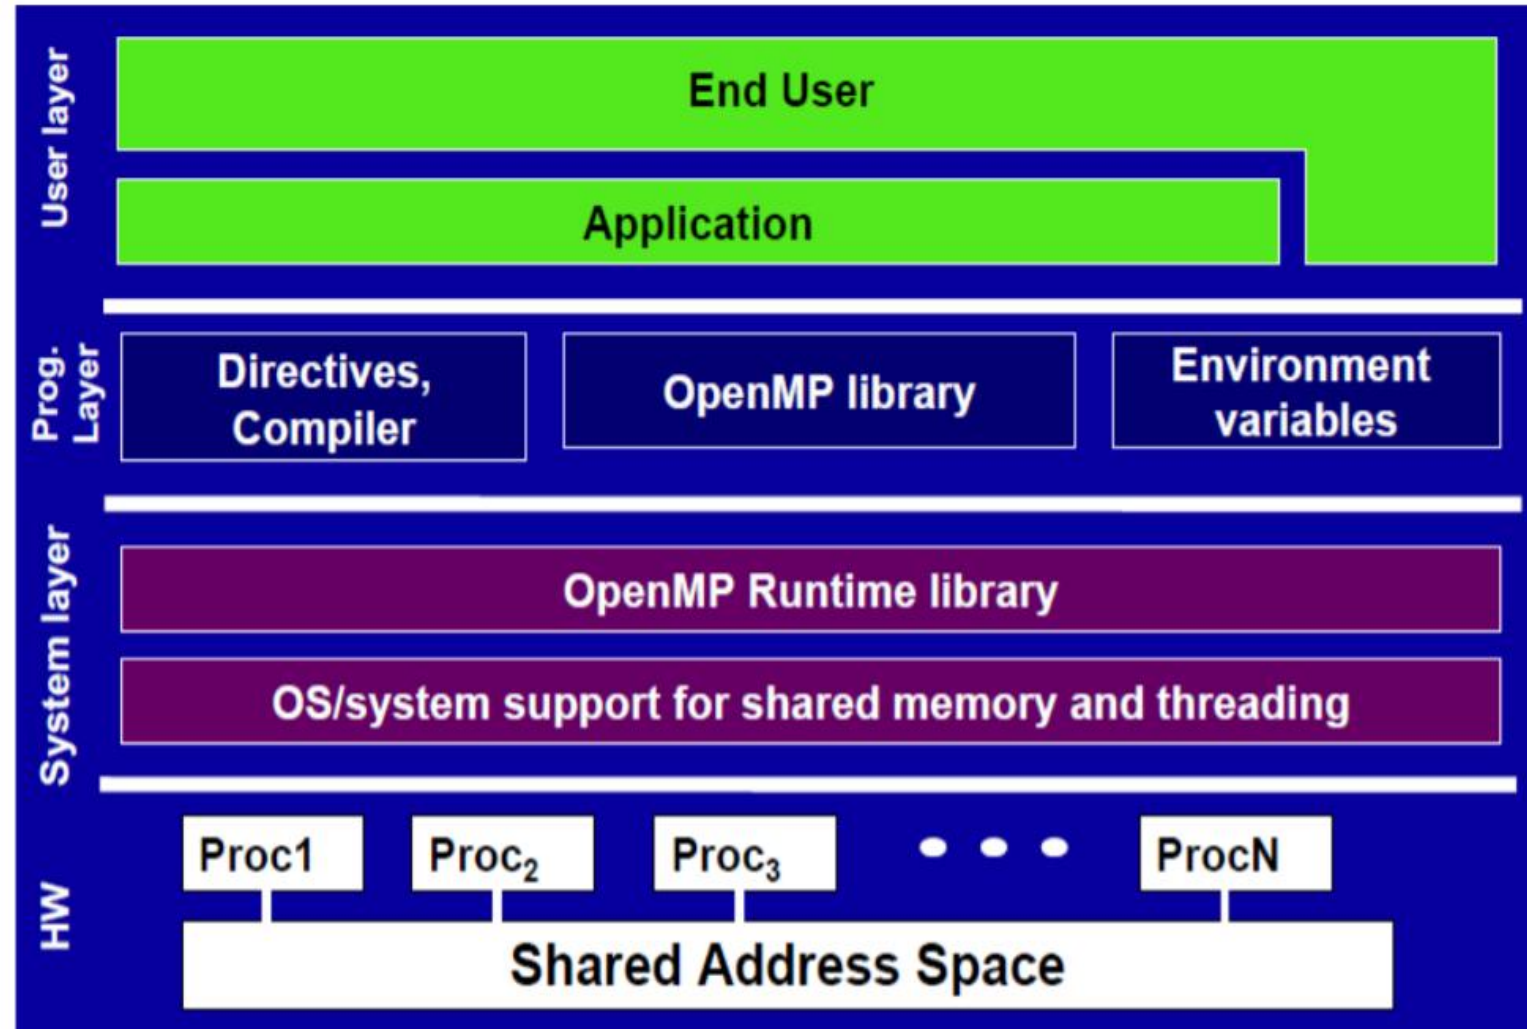
- OpenMP does not specify parallel I/O.
- It is up to the programmer to ensure that I/O is conducted correctly within the context of a multithreaded program.

- Memory Model

- Threads can “cache” their data and are not required to maintain exact consistency with real memory all the time.
- When it is critical that all threads view a shared variable identically, the programmer is responsible for ensuring that the variable is updated by all threads as needed. (*flush*)


1. OpenMP : API

- Architecture



1. OpenMP : API

- OpenMP Language Extensions



2. OpenMP Programming

```
%%Program hello.c
#include <stdio.h>
#include <omp.h>

int main(void)
{
    #pragma omp parallel
    printf("Hello, world.\n");
    return 0;
}
```

Compiling the program

```
$ gcc -fopenmp hello.c -o hello
```

Output

Hello, world.

Hello, world.

2. OpenMP Programming

- Directives

- An OpenMP executable directive applies to the succeeding structured block or an OpenMP Construct. A “structured block” is a single statement or a compound statement with a single entry at the top and a single exit at the bottom.

- Clauses

- Not all the clauses are valid on all directives. The set of clauses that is valid on a particular directive is described with the directive. Most of the clauses accept a comma-separated list of list items. All list items appearing in a clause must be visible.

- Runtime Library Routines

- Execution environment routines affect and monitor threads, processors, and the parallel environment. Lock routines support synchronization with OpenMP locks. Timing routines support a portable wall clock timer. Prototypes for the runtime library routines are defined in the file “omp.h”.

2. OpenMP Programming : *Parallel*

Directives

- Parallel
- For
- Sections
- Single
- Task
- Master
- Critical
- Barrier
- Taskwait
- Atomic
- Flush
- Ordered
- Threadprivate

Clauses

- Default
(shared/none)
- Shared
- Private
- Firstprivate
- Lastprivate
- Reduction
- Copyin
- copyprivate

Run time variables

- omp_set_num_threads
- omp_get_num_threads
- omp_get_max_threads
- omp_get_thread_num(void)
-etc

2. OpenMP Programming

- Directives

- An OpenMP executable directive applies to the succeeding structured block or an OpenMP Construct. A “structured block” is a single statement or a compound statement with a single entry at the top and a single exit at the bottom.
- They are case sensitive
- Starts with `#pragma omp`
- Directives cannot be embedded in embedded within continued statements, and statements cannot be embedded within directives.
- Only one directive-name can be specified per directive

2. OpenMP Programming: Directives

The parallel construct forms a team of threads and starts parallel execution.

```
#pragma omp parallel [clause[,]clause...] new-line
```

Structured-block

Clause: *if(scalar-expression)*

num_threads(integer-expression)

default(shared/none)

private(list)

firstprivate(list)

shared(list)

copyin(list)

reduction(operator:list)

Restrictions

- A program which branches into or out of a parallel region is non-conforming.
- A program must not depend on any ordering of the evaluations of the clauses of the parallel directive, or on any side effects of the evaluations of the clauses.
- At most one ***if*** clause can appear on the directive.
- At most one ***num_threads*** clause can appear on the directive. The ***num_threads*** expression must evaluate to a positive integer value.

2. OpenMP Programming: Directives

```
#pragma omp parallel [clause[,]clause...]  
new-line
```

Structured-block

Clause: if (*scalar-expression*)

 num_threads(*integer-expression*)

 default(*shared/none*)

 private(*list*)

 firstprivate(*list*)

 shared(*list*)

 copyin(*list*)

 reduction(*operator:list*)

- A team of threads is created to execute the parallel region
- A thread which encounters the parallel construct becomes master thread
- The thread id of master is 0
- All threads including master thread executes parallel region
- **omp_get_thread_num()** provides thread id
- There is implied barrier at the end of a parallel region.
- If a thread in a team executing a parallel region encounters another parallel directive, it creates a new team, and that thread is master of new team.

2. OpenMP Programming: Directives

#pragma omp parallel [*clause*[,]*clause*...]
new-line

Structured-block

Clause: **if**(*scalar-expression*)

num_threads(*integer-expression*)

default(*shared/none*)

private(*list*)

firstprivate(*list*)

shared(*list*)

copyin(*list*)

reduction(*operator:list*)

- If execution of a thread terminates while inside a parallel region, execution of all threads in all teams terminates.
- The order of termination of threads is unspecified
- All the work done by a team prior to any barrier which the team has passed in the program is guaranteed to be complete.
- The amount of work done by each thread after the last barrier that it passed and before it terminates is unspecified

2. OpenMP Programming: Clauses

```
#pragma omp parallel [clause[,clause...]
```

new-line

Structured-block

Clause: **if**(*scalar-expression*)

num_threads(*integer-expression*)

default(*shared/none*)

private(*list*)

firstprivate(*list*)

shared(*list*)

copyin(*list*)

reduction(*operator:list*)

If Clause

- A structured block is executed in parallel if the evaluation of **if()** clause is evaluated as true.
- A missing if clause is equivalent to an if clause that evaluates true.
- At most one **if ()** clause can appear on the directive.

2. OpenMP Programming : if clause

```
%%Program hello.c
```

```
%% if () clause
```

```
#include <stdio.h>
```

```
#include <omp.h>
```

```
int main(void)
```

```
{
```

```
int par=0;
```

```
#pragma omp parallel if(par==1) num_threads(4)
```

```
    printf("Hello, world.\n");
```

```
    return 0;
```

```
}
```

What is the result of the program?

Hello, world.

OR

Hello, world.

Hello, world.

Hello, world.

Hello, world.

2. OpenMP Programming: clauses num_threads

#pragma omp parallel [*clause*[,]*clause*...] *new-line*

Structured-block

Clause: **if**(*scalar-expression*)

num_threads(*integer-expression*)

default(*shared/none*)

private(*list*)

firstprivate(*list*)

shared(*list*)

copyin(*list*)

reduction(*operator:list*)

num_threads () Clause

- **num_threads** must evaluate to positive integer.
- It sets number of threads for the execution of parallel region
- Some of the execution environment routines
 - To set number of threads
void omp_set_num_threads(int num_threads);
 - To get number of threads
int omp_get_num_threads(void);
 - To find number of threads which can for a team
int omp_get_max_threads(void);
 - To get thread id
int omp_get_thread_num(void);

2. OpenMP Programming : num_threads clause

```
%%Program hello.c
```

```
%% if () clause
```

```
#include <stdio.h>
```

```
#include <omp.h>
```

```
int main(void)
```

```
{
```

```
int par=0;
```

```
#pragma omp parallel if(par==1) num_threads(6)
```

```
    printf("Hello, world.\n");
```

```
    return 0;
```

```
}
```

What is the result of the program?

Hello, world.

OR

Hello, world.

Hello, world.

Hello, world.

Hello, world.

Hello, world.

Hello, world.

2. OpenMP Programming: num_threads

```
%%Program hello.c
```

```
%% if () clause
```

```
#include <stdio.h>
```

```
#include <omp.h>
```

```
int main(void)
```

```
{
```

```
int par=1;
```

```
#pragma omp parallel if(par==1) num_threads(6)
```

```
    printf("Hello, world.\n");
```

```
    return 0;
```

```
}
```

What is the result of the program?

Hello, world.

OR

Hello, world.

Hello, world.

Hello, world.

Hello, world.

Hello, world.

Hello, world.

2. OpenMP Programming: Clauses

`#pragma omp parallel [clause[,]clause...] new-line`

Structured-block

Clause: `if(scalar-expression)`

`num_threads(integer-expression)`

`default(shared/none)`

`private(list)`

`firstprivate(list)`

`shared(list)`

`copyin(list)`

`reduction(operator:list)`

Default(shared/none) , shared(list) Clause

- **default (shared)** clause causes all variables referenced in the construct which have implicitly determined sharing attributes to be shared.
- **default(none)** clause requires that each variable which is referenced in the construct, and that does not have a predetermined sharing attribute, must have its sharing attribute explicitly determined by being listed in a data sharing attribute clause
- **shared(list)** : One or more list items must be shared among all the threads in a team.

2. OpenMP Programming: Directives

`#pragma omp parallel [clause[,clause...] new-line`

Structured-block

Clause: `if(scalar-expression)`

`num_threads(integer-expression)`

`default(shared/none)`

`private(list)`

`firstprivate(list)`

`shared(list)`

`copyin(list)`

`reduction(operator:list)`

private (list) Clause

- **private (list)** clause declares one or more list items must be private to a thread.
- A list item that appears in the **reduction** clause of a parallel construct must not appear in a **private** clause on a work-sharing construct.

2. OpenMP Programming: Clauses – default(shared)


```
// N: total number of iterations
#pragma omp parallel default(shared)
private(threadnum, numthreads, low, high, i)
{
    int threadnum = omp_get_thread_num(),
        numthreads = omp_get_num_threads();

    int low = N*threadnum/numthreads,
        high = N*(threadnum+1)/numthreads;

    for (i=low; i<high; i++)
        a[i]=b[i]+c[i]
}
```

Shared : A, B, C

Other variables are default



2. OpenMP Programming: Directives – Default(none)


```
// N: total number of iterations
#pragma omp parallel default(none)
shared(A,B,C,N) private(threadnum, numthreads,
low,high, i)
{
    int threadnum = omp_get_thread_num(),
        numthreads = omp_get_num_threads();

    int low = N*threadnum/numthreads,
        high = N*(threadnum+1)/numthreads;

    for (i=low; i<high; i++)
        a[i]=b[i]+c[i]
}
```

Shared : A, B, C

Other variables are default



2. OpenMP Programming: Directives - shared

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int main (void) {
    int x=0;
    #pragma omp parallel shared(x)
    {
        int tid=omp_get_thread_num();
        x=x+1;
        printf("Thread [%d] value of x is %d \n",tid,x);
    }

    return 0;
}
```

```
Thread [1] value of x is 1
Thread [0] value of x is 2
Thread [2] value of x is 3
Thread [3] value of x is 4
```

```
Thread [0] value of x is 2
Thread [1] value of x is 1
Thread [2] value of x is 3
Thread [3] value of x is 4
```

Since the x is shared, the change in one thread is visible to all other threads too.

2. OpenMP Programming: Directives - private

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int main (void) {
    int x=0;
    printf("x value ouside parallel:%d\n",x);
    #pragma omp parallel private(x)
    {
        x=10;
        int tid=omp_get_thread_num();
        x=x+1;
        printf("Thread [%d] value of x is %d \n",tid,x);
    }

    return 0;
}
```

```
x value ouside parallel:0
Thread [0] value of x is 11
Thread [2] value of x is 11
Thread [1] value of x is 11
Thread [3] value of x is 11
```

The variable `x` is private here . So the change of value performed in one thread is not visible to other threads.

2. OpenMP Programming: Directives - Private

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
int main (void) {
    int x=10, tid;
    printf("x value ouside parallel:%d\n",x);
    #pragma omp parallel num_threads(4) private(x) private(tid)
    {

        int tid=omp_get_thread_num();
        printf("\n 1. Thread [%d] value of x is %d \n",tid,x);
        x=15;
        printf("\n 2. Thread [%d] value of x is %d \n",tid,x);
        x=x+1;
        printf("\n 3. Thread [%d] value of x is %d \n",tid,x);
    }
    return 0;
}
```

```
x value ouside parallel:10
1. Thread [1] value of x is 6683504
2. Thread [1] value of x is 15
3. Thread [1] value of x is 16
1. Thread [3] value of x is 17143736
2. Thread [3] value of x is 15
3. Thread [3] value of x is 16
1. Thread [2] value of x is 6683600
2. Thread [2] value of x is 15
3. Thread [2] value of x is 16
1. Thread [0] value of x is 0
2. Thread [0] value of x is 15
3. Thread [0] value of x is 16
```

The variable x is private here. The value of x is 10 before parallel region. Since x is private, the value 10 is not reflected in threads. When the parallel region is entered, the x contains some garbage value as it is not initialed .

2. OpenMP Programming: Directives - Shared

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
int main (void) {
    int x=10, tid;
    printf("x value ouside parallel:%d\n",x);
    #pragma omp parallel num_threads(4) shared(x) private(tid)
    {
        int tid=omp_get_thread_num();
        printf("\n 1. Thread [%d] value of x is %d \n",tid,x);
        x=15;
        printf("\n 2. Thread [%d] value of x is %d \n",tid,x);
        x=x+1;
        printf("\n 3. Thread [%d] value of x is %d \n",tid,x);
    }
    return 0;
}
```

```
x value ouside parallel:10
1. Thread [2] value of x is 10
2. Thread [2] value of x is 15
3. Thread [2] value of x is 16
1. Thread [3] value of x is 10
2. Thread [3] value of x is 15
3. Thread [3] value of x is 16
1. Thread [1] value of x is 10
2. Thread [1] value of x is 15
3. Thread [1] value of x is 16
1. Thread [0] value of x is 10
2. Thread [0] value of x is 15
3. Thread [0] value of x is 16
```

X is shared. X is assigned value 15 inside parallel region.
Each thread is assigning value as 15. And updating it with $x=x+1$; But update is not getting reflected in other threads.

2. OpenMP Programming: Directives - Shared

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
int main (void) {
    int x=10, tid;
    printf("x value outside parallel:%d\n",x);
    #pragma omp parallel num_threads(4) shared(x) private(tid)
    {

        int tid=omp_get_thread_num();
        printf("\n 1. Thread [%d] value of x is %d \n",tid,x);
        //x=15;

        printf("\n 2. Thread [%d] value of x is %d \n",tid,x);
        x=x+1;

        printf("\n 3. Thread [%d] value of x is %d \n",tid,x);
    }
    return 0;
}
```

```
x value outside parallel:10
1. Thread [1] value of x is 10
2. Thread [1] value of x is 10
3. Thread [1] value of x is 11
1. Thread [3] value of x is 10
2. Thread [3] value of x is 11
3. Thread [3] value of x is 12
1. Thread [0] value of x is 10
2. Thread [0] value of x is 12
3. Thread [0] value of x is 13
1. Thread [2] value of x is 10
2. Thread [2] value of x is 13
3. Thread [2] value of x is 14
```

x is shared. No assignment in the parallel region.
Update is getting reflected in all other threads.
Synchronization is job of programmer.

2. OpenMP Programming: Directives

#pragma omp parallel [*clause*[,]*clause*...] *new-line*

Structured-block

Clause: **if**(*scalar-expression*)

num_threads(*integer-expression*)

default(*shared*/*none*)

private(*list*)

firstprivate(*list*)

shared(*list*)

copyin(*list*)

reduction(*operator*:*list*)

firstprivate (list) Clause

- **firstprivate (list)** clause declares one or more list items to be private to a thread and initializes each of them with that the corresponding original item has when the construct is encountered.
- For a **firstprivate** clause on a **parallel** construct, the initial value of the new item is the value of the original list item that exists immediately prior to the **parallel** construct for the thread that encounters the construct.

2. OpenMP Programming: Directives - Private

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
int main (void) {
    int x=10, tid;
    printf("x value ouside parallel:%d\n",x);
    #pragma omp parallel num_threads(4) private(x) private(tid)
    {

        int tid=omp_get_thread_num();
        printf("\n 1. Thread [%d] value of x is %d \n",tid,x);
        x=15;
        printf("\n 2. Thread [%d] value of x is %d \n",tid,x);
        x=x+1;
        printf("\n 3. Thread [%d] value of x is %d \n",tid,x);
    }
    return 0;
}
```

x value ouside parallel:10

```
1. Thread [1] value of x is 6683504
2. Thread [1] value of x is 15
3. Thread [1] value of x is 16
1. Thread [3] value of x is 17143736
2. Thread [3] value of x is 15
3. Thread [3] value of x is 16
1. Thread [2] value of x is 6683600
2. Thread [2] value of x is 15
3. Thread [2] value of x is 16
1. Thread [0] value of x is 0
2. Thread [0] value of x is 15
3. Thread [0] value of x is 16
```

X is private. While entering parallel region the x gets assigned with some value.

2. OpenMP Programming: Directives - Private

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
int main (void) {
    int x=10, tid;
    printf("x value ouside parallel:%d\n",x);
    #pragma omp parallel num_threads(4) private(x) private(tid)
    {

        int tid=omp_get_thread_num();
        printf("\n 1. Thread [%d] value of x is %d \n",tid,x);

        x=x+1;

        printf("\n 3. Thread [%d] value of x is %d \n",tid,x);
    }
    return 0;
}
```

x value ouside parallel:10

```
1. Thread [2] value of x is 6683600
3. Thread [2] value of x is 6683601
1. Thread [0] value of x is 0
3. Thread [0] value of x is 1
1. Thread [1] value of x is 6683504
3. Thread [1] value of x is 6683505
1. Thread [3] value of x is 14195704
3. Thread [3] value of x is 14195705
```

Shared x. x is already with some garbage value. X is private to each thread.

2. OpenMP Programming: Directives - Private

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
int main (void) {
    int x=10, tid;
    printf("x value ouside parallel:%d\n",x);
    #pragma omp parallel num_threads(4) firstprivate(x) private(tid)
    {

        int tid=omp_get_thread_num();
        printf("\n 1. Thread [%d] value of x is %d \n",tid,x);

        x=x+1;

        printf("\n 3. Thread [%d] value of x is %d \n",tid,x);
    }
    return 0;
}
```

```
x value ouside parallel:10

1. Thread [2] value of x is 10
3. Thread [2] value of x is 11
1. Thread [1] value of x is 10
3. Thread [1] value of x is 11
1. Thread [3] value of x is 10
3. Thread [3] value of x is 11
1. Thread [0] value of x is 10
3. Thread [0] value of x is 11
```

First private works as follows.

1. assign the value as in main thread before parallel region.
2. thread is private once it enters to parallel region.

2. OpenMP Programming: Directives

A loop Construct specifies that the iterations of loops will be distributed among and executed by the encountering team of threads.

```
#pragma omp for [clause[,]clause...] new-line  
for-loops
```

Clause: **private**(*list*)

firstprivate(*list*)

lastprivate(*list*)

reduction(*operator:list*)

schedule(*kind[,chunk_size]*)

collapse(*n*)

ordered

nowait

Example

```
#pragma omp parallel  
  #pragma omp for  
  for (i=0; i<N; i++) {  
    // do something with i  
  }
```

Index

- OpenMP
 - Program Structure
 - Directives : *parallel* , *for*
 - Clauses
 - If
 - Private
 - Shared
 - Default(shared/none)
 - Firstprivate
 - Num_threads

Reference

Text Books and/or Reference Books:

1. Professional CUDA C Programming – John Cheng, Max Grossman, Ty McKercher, 2014
2. B.Wilkinson, M.Allen, "Parallel Programming: Techniques and Applications Using Networked Workstations and Parallel Computers", Pearson Education, 1999
3. I.Foster, "Designing and building parallel programs", 2003
4. Parallel Programming in C using OpenMP and MPI – Micheal J Quinn, 2004
5. Introduction to Parallel Programming – Peter S Pacheco, Morgan Kaufmann Publishers, 2011
6. Advanced Computer Architectures: A design approach, Dezso Sima, Terence Fountain, Peter Kacsuk, 2002
7. Parallel Computer Architecture : A hardware/Software Approach, David E Culler, Jaswinder Pal Singh Anoop Gupta, 2011
8. Introduction to Parallel Computing, Ananth Grama, Anshul Gupta, George Karypis, Vipin Kumar, Pearson, 2011

Reference

Acknowledgements

1. Introduction to OpenMP <https://www3.nd.edu/~z xu2/acms60212-40212/Lec-12-OpenMP.pdf>
2. Introduction to parallel programming for shared memory Machines <https://www.youtube.com/watch?v=LL3TAHpxOig>
3. OpenMP Application Program Interface Version 2.5 May 2005
4. OpenMP Application Program Interface Version 5.0 November 2018



Thank You